



Traditional IT Ops to SRE Mindset

1



Introduction & Evolution

2



Real-World Story

3

The Night of the Big Incident

- It was **11:50 PM on a Monday**.
- The stock trading application for **National Financial Exchange** had just gone into production with a **minor configuration change — approved** through **Change Management**, documented in **ITIL, tested in UAT**.
- **Everything looked fine.**
- But at **9:15 AM next morning — right at market opening — the system slowed down.**
- **Within minutes:**
 - Trade orders started failing
 - Brokers couldn't log in
 - Transactions got stuck in queue
 - Twitter exploded: "Exchange Services Down!"
 - News channels picked it up.
 - In just 30 minutes — **market loss ₹100 Crores.**

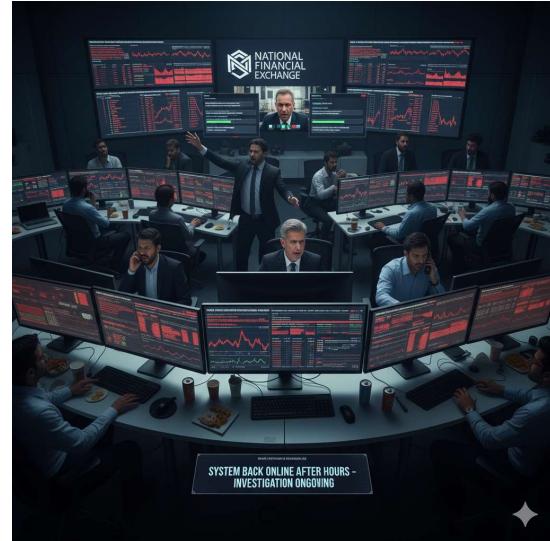


4



Ops Team — Fix

- The **Ops team rushed to fix it.**
 - They **checked logs, restarted servers, increased CPU.**
 - Senior management called every 10 minutes.
 - **Incident reports were created.**
 - **Root Cause Analysis meetings were scheduled.**
- Finally, after **2 hours**, the **system was back.**
- But the **damage was already done.**



5

A Week Later — The Real Problem Emerged

- In a high-level review meeting, the CIO didn't ask, "*Who made the mistake?*".
- He asked:
 - "Why are we still **depending** on people to **manually rescue systems?**"
 - "Why don't our **systems heal themselves?**"
 - "Why can't we **predict failures** instead of reacting?"
 - "Why do we find problems **after users complain?**"
- There was silence.
- Then, one Ops lead said something that changed everything: "**We have been doing Operations but not Engineering.**"



6

Beginning of the SRE Journey

- They hired **reliability engineers**.
- They **automated recovery**.
- They **eliminated 70% tickets** using **scripts**.
- They introduced **real-time observability** — not just CPU/Memory — but latency, transaction drops, log anomalies.
- Now, if the same issue happens?
 - The **system auto-detects** slow database connections
 - It **automatically shifts** to read-replica
 - No downtime.
 - No Twitter complaints.
 - No financial loss.



7

Discussion Questions

- **Question**
 - In the story, Person X was praised as a "**hero**" for fixing the incident manually.
 - In your experience, **what are the risks or hidden costs** associated with a team constantly relying on "heroes" for production stability?
- **Question**
 - Can you recall a specific instance where your team's expertise was spent on a fix that you knew—even as you were doing it—was only **temporary**?
 - What was the **repetitive task** you wished you could automate?
- **Question**
 - In the story, incident took 45 minutes to resolve.
 - Where was most of that time spent: **Alerting, Diagnosis, or Mitigation?**



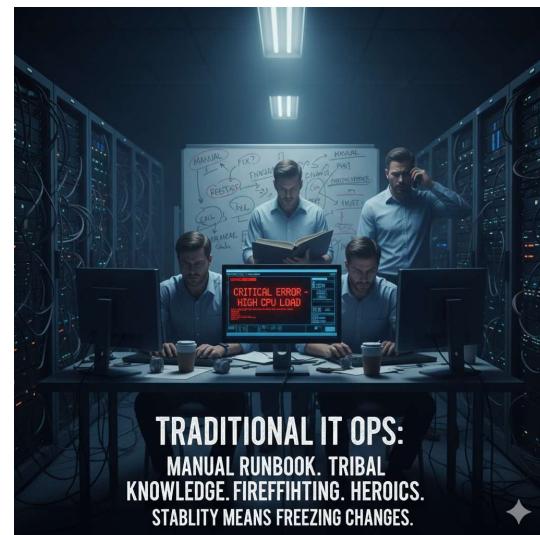
8



9

What is IT Operations

- IT Operations **refers to all the activities** required to **keep** an organization's **IT infrastructure** and **applications** **running smoothly, securely, and efficiently.**
- It involves **managing hardware, servers, networks, databases, security, backups, help desk, monitoring, incident handling, and compliance.**
- **ITOps focuses on the day-to-day execution of tasks** necessary to deliver and support IT services to both **internal users** (employees) and **external customers**



10

IT Operations ... continue

- **Traditional IT Operations** refers to the **legacy approach of managing IT systems**, where teams focus on:
 - **Maintaining** servers, networks, storage
 - **Handling** incidents through **tickets**
 - **Monitoring** systems manually
 - **Managing** data centers and on-prem infrastructure
 - **Following** ITIL-based processes like Change, Incident, and Problem Management
- **Goal in Traditional Ops:**
 - **Keep systems running**, stable, and secure — but **using manual processes** and **operational control**.
- The model is characterized by a deep **siloing** of responsibilities:
 - **Development (Dev)**
 - Writes code and focuses on new features.
 - **Operations (Ops)**
 - Manages the running system, infrastructure, security, and deployment.
 - **Handoffs**
 - Code is completed by Dev and then "thrown over the wall" to Ops for deployment and maintenance.



11

Key Responsibilities of IT Operations

Area	Description
Infrastructure Management	Setup & manage servers, networks, storage
Monitoring & Maintenance	Track performance, uptime, logs, health
Deployment & Change Control	Deploy software with approvals
Incident & Problem Management	Respond to system failures and perform root cause analysis
Security & Compliance	Protect from threats, control access, audit
Backup & Disaster Recovery	Data backup, replication, recovery plans
Help Desk / Service Desk	User support, ticket resolution



12

Types of ITOps Teams and Function

Type of Team/Function	Core Focus Area	Mindset & Goal
Service Desk/Help Desk	First-line support and incident logging; handling employee issues (password resets, device failures).	Customer Satisfaction: Act as the single point of contact (SPOC) and resolve issues quickly.
Infrastructure Management	Physical and virtual hardware (servers, storage, data centers, cloud resources), operating systems, and network components.	Availability: Ensure the underlying platforms are always running and scalable.
Network Operations Center (NOC)	Real-time monitoring of the network, applications, and services; identifying and triaging high-priority incidents.	Proactive Alerting: Minimize downtime by detecting and escalating problems before they cause outages.
Security Operations (SecOps)	Implementing security controls, monitoring for threats , managing firewalls, and responding to security incidents.	Data Protection: Maintain system integrity and compliance with security policies.
Data & Storage Operations	Backups , disaster recovery, data retention, and ensuring the performance and availability of database servers.	Resiliency: Guarantee data integrity and the ability to recover service after a catastrophic failure.
Application Operations	Managing the health, deployment, and performance tuning of specific business applications once they are in production.	Performance: Ensure critical business applications run smoothly for end-users. ⁴



13

Approach and Processes

- The Traditional Ops approach is characterized by a "**throw it over the wall**" handoff and a focus on minimizing change to maximize stability.
 - The Handoff ("Throw It Over the Wall")**
 - Once the **Development (Dev)** team **completes** the **code** and documentation, it is **handed over** to the **Operations (Ops)** team for deployment and ongoing maintenance.
 - This **often leads** to the infamous "**it works on my machine**" problem.
 - Change Management Process**
 - Changes** (*code deployments, configuration updates*) are viewed as the **primary source of risk**.
 - They are **managed** through slow, multi-layered approval processes, often involving a **Change Advisory Board (CAB)**.
 - The **goal** is to **restrict** and minimize the **frequency of changes**.
 - Reactionary Incident Response**
 - Incidents** are typically **resolved** by engineers manually **logging into systems**, diagnosing the failure, and applying a fix.
 - The **primary focus** is on **restoring service immediately**, often leaving the underlying cause unaddressed or documented only in a tribal knowledge format.
 - Monitoring Focus**
 - Monitoring** often focuses on basic **system health metrics** (*CPU, memory, disk usage*) rather than the application's actual business service health.
 - Alerts** can be numerous and **noisy**, leading to **alert fatigue**.



14

ITOps Workflow

- Scenario: Deploying a New Application
 1. **Development**
 - The Dev team finishes the application code and hands off the build files and a long installation manual (the **handoff**).
 2. **Request**
 - The Ops manager receives a **Change Request (CR)** ticket to deploy the application.
 3. **Approval**
 - The CR goes to the **Change Advisory Board (CAB)**, which meets once a week.
 - The CR is scrutinized for potential impact and approved after three days.
 4. **Deployment (The Toil)**
 - An **Ops engineer manually** follows the installation manual.
 - They **log** into the **production** server via SSH, **manually configure** the web server, **install** dependencies, copy the files, and **restart** the service.
 - This is often **done during a scheduled maintenance window** at 2 AM.
 5. **Monitoring**
 - The Ops team sets up basic monitoring to check if the server is alive (ping) and the CPU/memory usage, satisfying the SLA.
- This model prioritizes **control and stability** over **speed and collaboration**, which is the primary driver for the evolution toward modern practices like DevOps and SRE.



15

ITOps | Benefits

- The **focus** on **control** and **process** was highly effective for the pre-internet era **when systems were static** and **changes were infrequent**.
 - **Stability & Control**
 - Strict, formal **processes** (like the *Change Advisory Board*, or CAB) **minimize changes**, leading to high stability for static, mission-critical systems.
 - **Clear Accountability**
 - Roles are **strictly defined** (*Network Admin*, *Database Admin*, *Sysadmin*), ensuring **clear ownership** and accountability for specific infrastructure components.
 - **Security & Compliance**
 - The separation of duties and rigid, documented procedures make it easier to meet **stringent regulatory requirements**, especially in finance or healthcare.
 - **Cost Predictability**
 - Operating on owned, on-premise hardware allows for **predictable, long-term Capital Expenditure** (*CapEx*) budgeting, rather than variable monthly cloud costs.



16

ITOps | Drawbacks

- The reliance on manual effort and **rigid processes creates friction** and **bottlenecks** in the modern business environment.
 - Slow-Release Cycles**
 - The sequential **Waterfall model** and formal **change management processes (CAB meetings)** mean **releases** are **infrequent, large, and high-risk ("big bang" deployments)**.
 - Adversarial Culture (Dev vs. Ops)**
 - The **siloed structure creates conflict**.
 - Developers prioritize speed; Operators prioritize stability**.
 - This results in mutual frustration and a **lack of shared ownership**.
 - High Operational Toil**
 - Engineers spend** a vast amount of time on **repetitive, manual, and error-prone tasks** (like *manual patching, server configuration, or log checks*).
 - This **toil prevents** them from focusing on strategic improvements.
 - Poor Incident Response**
 - Troubleshooting** often involves **manually logging** into servers and **following runbooks**, leading to long **Mean Time To Resolution (MTTR)** and costly downtime.



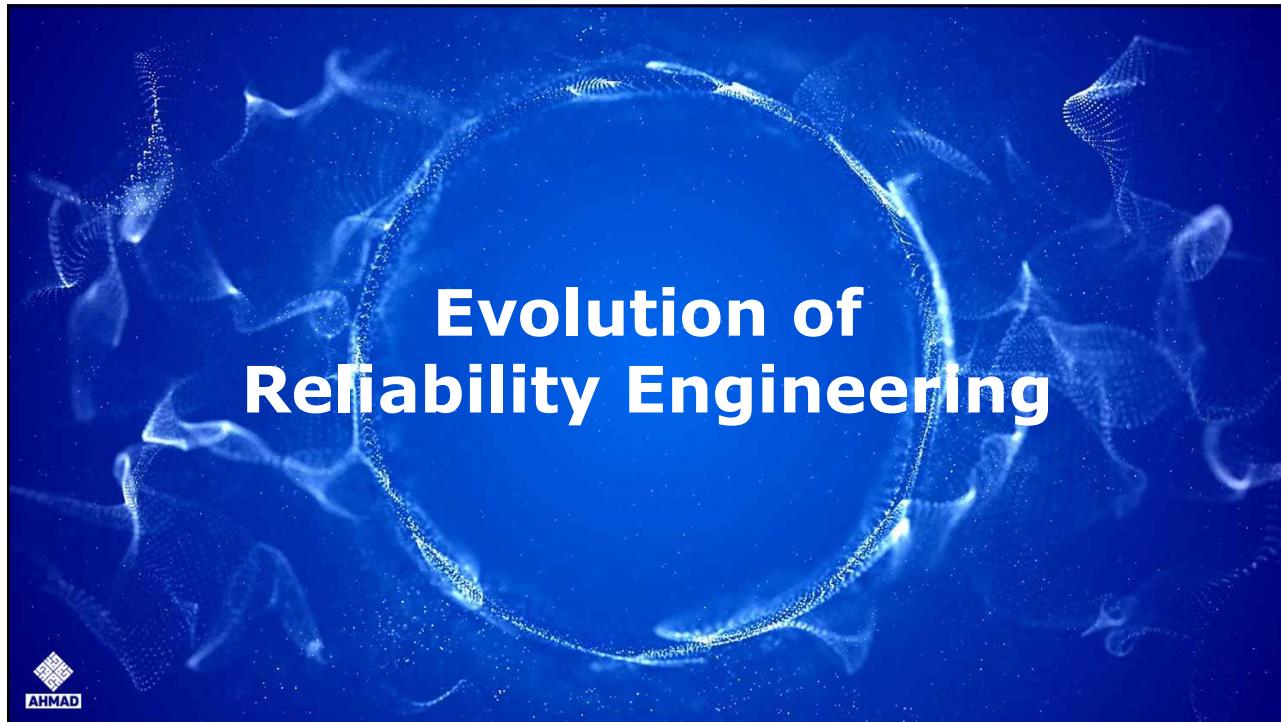
17

Why Traditional ITOps Fails in New IT

Traditional IT Ops	Modern Digital IT	Why Alignment Fails
Manual, Ticket-Based	Automated, Self-Healing	<ul style="list-style-type: none"> Manual ticket handling is slow, error-prone, and dependent on human availability. Modern systems need real-time automated responses, self-healing scripts, and zero-touch resolutions to maintain uptime.
Reactive Issue Handling	Predictive and Preventive Monitoring	<ul style="list-style-type: none"> Traditional ops waits for issues to happen (reactive). Modern systems predict failures using metrics, anomaly detection, and AI — fixing issues before customers are impacted.
On-Prem Infrastructure	Cloud, Containers, Serverless	<ul style="list-style-type: none"> Static, hardware-heavy infrastructure cannot support rapid scaling, multi-region deployment, or cloud-native architectures required by digital platforms.
Monthly Deployments	Multiple Deployments per Day	<ul style="list-style-type: none"> Legacy IT relies on approval boards and long change cycles. Digital businesses (Amazon, Netflix) deploy continuously using CI/CD, feature flags, and automated testing.
Static Capacity Planning	Dynamic Auto-Scaling	<ul style="list-style-type: none"> Traditional IT purchases capacity for peak load, leading to high cost and underutilization. Digital IT auto-scales instantly based on real-time demand in cloud environments.
Human Monitoring	AI/Analytics-based Observability	<ul style="list-style-type: none"> Manual monitoring (logs, dashboards) misses hidden performance issues. Modern monitoring uses AI, distributed tracing, and real-time analytics to provide deep visibility.



18



19

What is Site Reliability Engineering

- **SRE (Site Reliability Engineering)** is an engineering approach introduced by **Google** to ensure **highly reliable, scalable, and efficient software systems** using **software engineering practices applied to operations**.
- The concept was pioneered at **Google in 2003** by Ben, who defined SRE as "**what happens when you ask software engineers to design an operations function**"
- It combines **software development + IT operations**, with a strong focus on **automation, reliability, performance, monitoring, and incident response**.
- "**Use software engineering to solve operational problems.**" Instead of managing operations manually (like traditional Ops), SRE automates them through code.



20

What is Reliability

- Reliability in Site Reliability Engineering (SRE) is the **single most important metric** and the **core focus** of the entire discipline.
- In simple terms, SRE views reliability as **how often and how well a service works as expected by its users**.
- It's **not just** about the **server being "up,"** but about the **user being able to successfully complete their task** (e.g., loading a webpage, sending a message, completing a purchase) quickly and without errors.



21

Reliability in Traditional IT vs SRE

Aspect	Traditional IT (Ops)	SRE Approach
Definition	Keep the server running	Ensure quality experience for users
Focus	System uptime	Overall user happiness (availability, speed, errors, security)
Measurement	Uptime (%) only	SLO, SLI, error rate, latency, scalability, resilience
Approach	Manual fixes, reactive	Automated, data-driven, proactive
Mindset	"Keep lights on"	"Engineer reliability like a product feature"



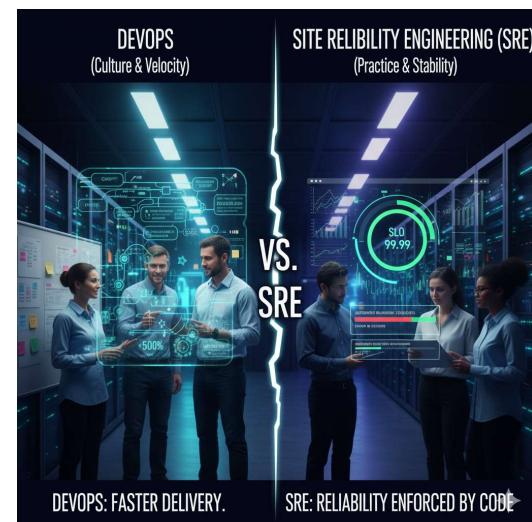
22



23

DevOps vs SRE

- While **DevOps** is a **philosophy** that **drives cultural change**, **SRE** is a prescriptive **practice** that defines **how to implement** the **reliability** and **automation** aspects of DevOps using software engineering methods.
- Both** DevOps and SRE **improve software delivery** and **operations**.
- But **they solve different problems**, **use different methods**, and **have different goals**.
- DevOps emphasizes** the **pipeline** (CI/CD), and **SRE emphasizes** production **reliability**. *Both complement each other* in modern organizations.



24

DevOps vs SRE | Key Differences

Aspect	DevOps (Philosophy/Culture)	SRE (Implementation/Role)
Goal	<ul style="list-style-type: none"> Increase the organization's ability to deliver applications and services at high velocity 	<ul style="list-style-type: none"> Define, achieve, and maintain specific levels of service reliability using quantifiable metrics
Focus	<ul style="list-style-type: none"> Collaboration and Speed. Breaking down organizational silos between Development and Operations 	<ul style="list-style-type: none"> Reliability and Automation. Applying software engineering to solve operations problems
Who does the work?	<ul style="list-style-type: none"> Dev + Ops teams together 	<ul style="list-style-type: none"> Specialized reliability engineers
Philosophy	<ul style="list-style-type: none"> "Developers and operations should collaborate" 	<ul style="list-style-type: none"> "Ops is a software engineering problem"
Metrics Focus	<ul style="list-style-type: none"> Deployment frequency, lead time, and mean time to recovery (MTTR) 	<ul style="list-style-type: none"> Service Level Objectives (SLOs), Error Budget consumption, and percentage of time spent on toil



25

DevOps vs SRE | Mapping

Statement	DevOps	SRE
How can we release faster ?	Yes	Yes
How can we stay reliable during fast releases?	Not fully defined	Core mission
Is it a culture ?	Yes	No — it's a role/discipline
Is it a team ?	No (Dev + Ops collaboration)	Yes — SRE team exists
Do they write automation code ?	Sometimes	Always



26

Pre-SRE Approaches



27

**What the World Did Before
Google Invented SRE**



28



29

Pre-SRE Approaches

- Before Site Reliability Engineering (SRE) emerged, organizations primarily used **Traditional IT Operations (IT Ops)** and **ITIL-based management** to run systems, maintain availability, and manage incidents.
- These approaches** were **process-heavy, manual, and reactive**.



PRE-SRE APPROACHES: MANUAL CONTROL,
SLOW CHANGES, AND THE AGE OF HEROICS



30

ITOps and System Administration

- Traditional IT Ops and Sysadmin were the **primary roles** responsible for **maintaining** production systems.
- **Focus:** The core focus was maintaining **stability** through manual control and **physical infrastructure management**.
- **Mindset**
 - The **environment** was **seen** as **static**. **Change** was **inherently risky**, and the **Sysadmin's job** was to be the **gatekeeper against failure**.
- **Key Activities**
 - **Server Maintenance**
 - Physically **racking** and **cabling** servers, **installing** and **patching** operating systems, and managing virtual machines **manually**.
 - **Network/Storage Management**
 - **Configuring** network **devices**, managing **firewalls**, and provisioning **storage arrays**.
 - **Scripting**
 - Using **basic shell scripts** (*Bash, PowerShell*) to **automate repetitive tasks** **but** primarily **focusing on execution rather than engineering**.



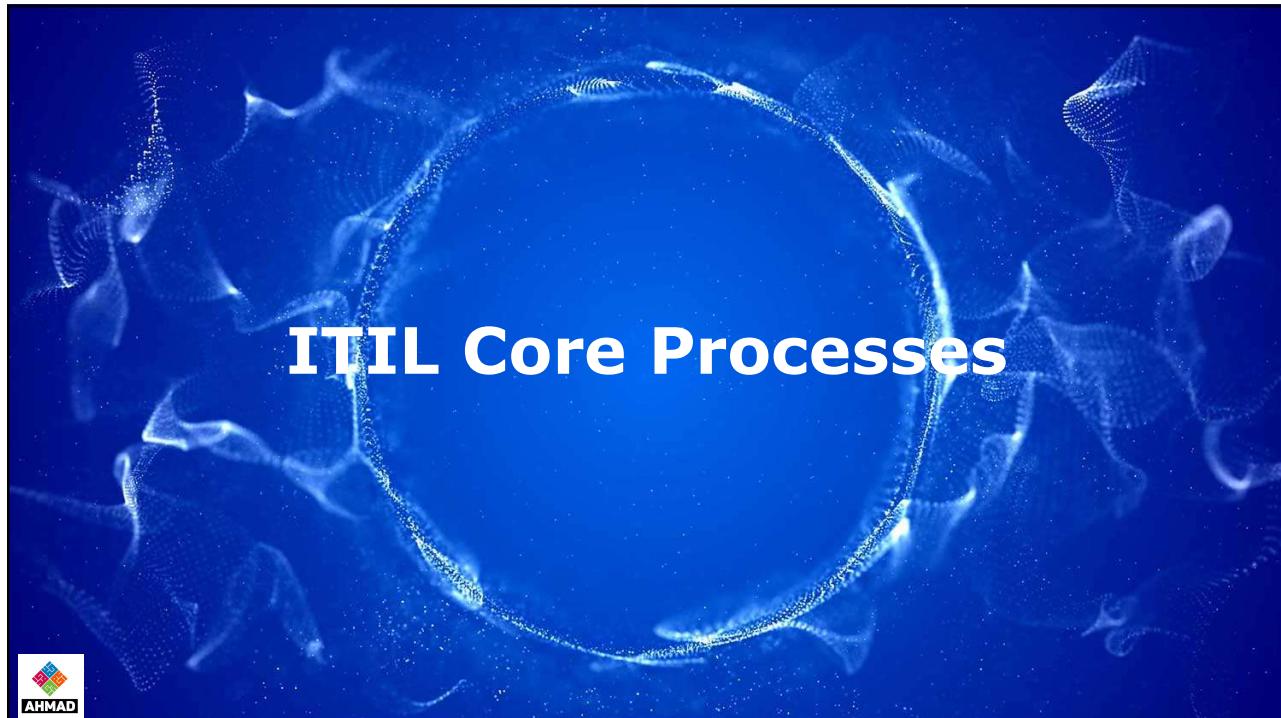
31

ITOps and Sysadmin ... continue

- **Characteristics:**
 - **Approach** → Manual, Reactive
 - **Focus** → Hardware, Servers, Uptime
 - **Work Type** → Tickets, Procedures, Checklists
 - **Automation** → Very limited
 - **Scaling** → Hard to scale
- **Example:**
 - An **e-commerce server goes down** during a sale.
 - **Admins manually log in, restart services, check logs**.
 - **Root cause analysis is done later**.
 - **No automated recovery** or self-healing.



32



33

ITIL-Based Management vs SRE Philosophy

- **ITIL** (Information Technology Infrastructure Library) is a comprehensive **framework** detailing **best practices** for **IT Service Management** (ITSM).
- It became the **dominant management structure** for **large, stable enterprises** in the **pre-SRE era**.
- Understanding **Mindset Difference**:

Mindset	ITIL	SRE
"How do we avoid failure completely?"	Yes	Not realistic; systems WILL fail
"How do we recover quickly with minimum impact?"	Not Primary focus	CORE IDEA of SRE

- **Key SRE Differentiator**
 - **SRE believes failure is normal** and focuses on **making systems resilient** to failure using automation, monitoring, and engineering practices.



34

ITIL-Based Management vs SRE Philosophy

Feature	ITIL-based Management (Traditional)	SRE Philosophy (Modern)
Core Goal	<ul style="list-style-type: none"> Control and Process Adherence. Focus on documented steps to maintain service quality. 	<ul style="list-style-type: none"> Reliability and Automation. Focus on achieving measurable outcomes (SLOs) through engineering.
View of Manual Work	<ul style="list-style-type: none"> Processes are mandatory. Manual execution of change forms and approvals is a required control point. 	<ul style="list-style-type: none"> Manual work is Toil that must be eliminated. If a process is repeated, it must be automated by code.
Approach to Risk	<ul style="list-style-type: none"> Risk Aversion. Minimize change to avoid incidents, leading to slow-release cycles. 	<ul style="list-style-type: none"> Risk Management. Define downtime allowance and use data to balance agility (speed) against reliability (stability).



35

ITIL Core Processes

- In a traditional, pre-SRE environment, the **three core ITIL processes** were handled with heavy human involvement:
 - Incident Management
 - Change Management
 - Problem Management



36

ITIL Core Processes ... continue

- **Incident Management**
 - **Legacy Approach**
 - Primarily **reactive** and **manual**.
 - **When a system failed:**
 - An **alert** (often generic) would page an on-call engineer.
 - The engineer would **manually log into servers**, check logs, and try to diagnose the issue based on their experience and paper/digital runbooks.
 - Mitigation involved **manual steps** like restarting services or failing over to a standby system.
 - **Result**
 - **Long Mean Time To Resolution** (MTTR), as human intervention and diagnosis are the slowest parts of the process.



37

ITIL Core Processes ... continue

- **Change Management**
 - **Legacy Approach**
 - ✓ Dominated by the **Change Advisory Board** (CAB).
 - ✓ A developer or operator wanting to deploy a change had to submit a detailed change request ticket.
 - ✓ The **CAB** (*a committee of senior stakeholders*) would **meet weekly** to manually **review** and **approve** or **deny** the change.
 - **Result**
 - ✓ **Releases were slow**, infrequent, and large ("big bang" deployments), which paradoxically made them high-risk, as one big change is harder to troubleshoot than many small ones.



38

ITIL Core Processes ... continue

- **Problem Management**

- **Legacy Approach**

- After an **incident** was resolved, a formal **Root Cause Analysis (RCA)** document was created.
- The focus was on **identifying the cause and proposing a solution** (e.g., "The server needs more memory").
- However, the **implementation** of that solution (*the actual fix*) often **got stuck** in the backlog or required another slow **CAB approval**.

- **Result**

- The **same problem** (and the resulting incident) often recurred because the operations team lacked the engineering time and authority to **implement permanent, automated fixes**, leading to cycles of recurring toil.



39

Core SRE Mindset



40

Core SRE Mindset

- SRE isn't just a role — it's a **mindset shift** in how organizations manage systems.
- It moves from **People + Manual Work** → to **Engineering + Automation**.
- The **core belief** is:
 - **Reliability** is a **feature**. It **must be designed, engineered, measured, and continuously improved — not just maintained** manually.
- At its heart, SRE is about **engineering reliability into systems** while still enabling innovation.
- The mindset is not "zero failure" but **controlled, measurable reliability** balanced with agility.



41

Reliability vs. Agility Trade-off

- In the modern IT world, companies want **rapid feature releases** (*agility*) and **high reliability** (*stability*) — but **doing both is difficult**.
- **Traditional Ops Mindset**
 - **Slow and safe change**
 - Stability achieved by **blocking frequent deployments**
 - **Focus on uptime**, not engineering
 - "Let's stop deployments to avoid failures."
- **SRE Mindset**
 - Reliability is not achieved by restricting changes — it is achieved by **engineering for resilience**
 - Use **automation, safe deployments, testing, observability**
 - Measure reliability scientifically
 - "We don't slow down releases; we make releases safer."



42

SRE in On-Prem and Hybrid Environments

- While **SRE principles** are **most famously** associated with **cloud-native** companies like Google, the **mindset** is **technology-agnostic** and **highly effective** in managing traditional **on-premise** and **hybrid** environments.
- The **fundamental** goal remains the **same: eliminate toil and manage risk through code**.

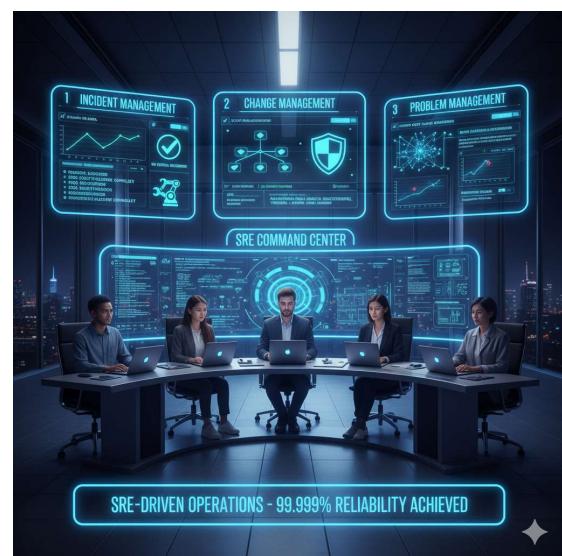
Principle	Cloud-Native SRE Approach	On-Prem/Hybrid SRE Approach
Infrastructure	Use Cloud Provider APIs to provision resources (<i>AWS CloudFormation, GCP Deployment Manager</i>).	Use Infrastructure as Code (IaC) tools like Terraform or Ansible to manage virtualization (VMware), network appliances, and bare-metal servers.
Self-Healing	Rely on container orchestrators (Kubernetes) and cloud-managed services for auto-scaling and health checks .	Build automation workflows using runbook automation engines or internal tools that script server restarts, failovers, and resource adjustments.
Challenge	The main challenge is overcoming the manual control culture ingrained in traditional data center management and securing the necessary permissions for automated tooling.	The main challenge is the inherent difficulty of automation in legacy systems that may lack modern APIs or rely on manual CLI interactions.



43

Management in the Post-SRE World

- In the SRE model, these traditional functions are transformed into engineering disciplines:
 - Change Management** to **Change Engineering**
 - Incident Management** to **Incident Response Engineering**
 - Problem Management** to **Reliability Engineering**



44

Management in the Post-SRE World

- **Change Management → Change Engineering**

- The primary goal of Change Management is still to ensure changes are deployed safely, but the mechanism is completely automated.
 - ✓ CAB Approvals → CI/CD automated approvals
 - ✓ Change Freeze → Feature flags & gradual rollouts
 - ✓ Night-time manual deployments → Auto deployments with rollback
 - ✓ Verbal approvals → GitOps-based approval control (*approvals are managed through Git repositories*)

- **Incident Management → Incident Response Engineering**

- The SRE approach aims to make incidents rare, reduce human involvement in mitigation, and ensure recovery is instantaneous.
 - ✓ Manual escalation (phone/email) → Automated alerting
 - ✓ No defined roles → Incident Command System
 - ✓ Focus: restore service manually → Focus: auto-remediation with scripts
 - ✓ Retrospective blame → Blameless postmortems



45

Management in the Post-SRE World

- **Problem Management → Reliability Engineering**

- The goal shifts from merely finding the root cause to guaranteeing that the problem cannot recur through code changes.
 - ✓ Root Cause Analysis (who did it?) → Blameless Postmortems (why did system fail?)
 - ✓ Document fixes only → Implement resilience improvements
 - ✓ No testing for resilience → Use reliability testing tool
 - ✓ Manual restoration → Self-healing architecture



46

Case Study



47

ITOps at a Stock Exchange



48

Stage

- Global **TradeX Stock Exchange** operated like most traditional financial institutions.
- It had a big data center, mainframe servers, and a Traditional IT Ops team responsible for monitoring systems, responding to tickets, and applying patches.
- Their operations followed:
 - **Manual Monitoring:** Teams watched dashboards 24/7 for alerts
 - **Ticket-Based Support:** Traders logged tickets when systems were slow
 - **Monthly Scheduled Maintenance:** Downtime allowed on weekends only
 - **Change Approvals via CAB:** Deployment approvals took 2-3 weeks
 - **Fix after Failure:** Teams waited for systems to crash before acting
- They believed: "**Stability means freezing changes.**" But the world around them was changing.



49

The Day Everything Broke

- One Monday morning, during peak trading hours (9:15 AM), the **market order execution system slowed down.**
- Trades that should complete in **40 milliseconds** were taking **4 seconds**. In stock trading, **4 seconds is a lifetime** — people lost money.
- The helpdesk opened **hundreds of tickets**.
- The **Ops Team Response:**
 - They manually restarted servers
 - Increased memory on application nodes
 - Filed incident reports to multiple teams
- **Total resolution time: 2 hours 20 minutes**
- **Result**
 - Millions lost in failed or delayed trades
 - Exchange reputation hit media and regulators
 - CEO called an emergency board meeting



50

Root Cause Analysis

- During the review, **key problems** surfaced:
 - **No early warning:** Manual monitoring didn't predict issue
 - **No automation:** Every fix required engineers and approvals
 - **No auto-scaling:** Load increased 300%, servers stayed same
 - **No real-time incident response:** Took 2 hours to coordinate teams
 - **Not enough observability:** Couldn't pinpoint the exact failure fast
- The **CIO said:** We don't need more **people watching screens**. We **need systems** that **heal themselves** before failing."



51

The Shift – From Ops to SRE

- They **introduced SRE** (Site Reliability Engineering) to transform operations:
 - **Ticket-based work → Automation-first (reduce tickets)**
 - SRE builds scripts & automation to **prevent issues before tickets are raised**
 - ✓ Auto-restart service, reallocate resources, auto-scale system etc.
 - **Stability via change blocking → Reliability through engineering**
 - Use engineering to **make changes safe, fast**, and reliable
 - ✓ Automated testing (CI/CD)
 - ✓ Canary Deployments
 - ✓ Blue-Green Releases
 - **Manual incident recovery → Self-healing systems**
 - Systems **auto-detect** and **auto-fix** failures without human involvement
 - ✓ High CPU usage → Auto-scale servers
 - ✓ Service crashes → Auto-restart container/VM
 - ✓ Region failure → Auto-failover to another region
 - ✓ Slow DB → Automatically add read replicas



52

The Shift – From Ops to SRE ... continue

- **General IT operations → Software engineers specialized in reliability**
 - Software engineers **who write** code to **solve operational problems**
 - ✓ Write Python/Go scripts for automation
 - ✓ Build Monitoring & Alerting systems
 - ✓ Create tools to manage capacity, failures, scaling
 - ✓ Use chaos engineering to test system resilience
- **Human monitoring → AI-based proactive observability**
 - **AI-driven monitoring, anomaly detection, tracing**
 - Use Observability platforms (*Prometheus, DataDog, Grafana, ELK*)
 - AI/ML detects anomalies before failure
 - Full-stack visibility: logs, metrics, traces, user experience
 - Business monitoring (*trade volume, revenue impact, failed orders*)



53

One Year Later: New World, New Results

Before (Traditional Ops)	After (SRE-based Ops)
4 major outages/year	0 major outages in last year
Incident Recovery Time: 2+ hours	Recovery Time: under 2 minutes (self-healing)
300+ monthly tickets	85% tickets eliminated via automation
Manual change approvals	Automated safe deployments (Canary, Blue/Green)
Monitoring only CPU & RAM	Deep Observability (Latency, traces, business metrics)



54

Q & A

Any concepts still unclear?

Thank you for attending

