



SRE Principles SLIs, SLOs & Error Budgets

1



Key SRE Principles

2

Core Principles of SRE

- SRE is defined by a set of **key principles** and **practices** that **govern** how **reliability** is **managed** and **measured**.
- Site Reliability Engineering (SRE) is about **making systems reliable using engineering practices**.
- Instead of fixing problems manually, SRE teams build automation, monitoring, and tools to keep services running smoothly.
- SRE is defined by a software engineering approach to operations, **centered on three core principles** that **drive all decisions**:
 - Embracing Risk and Setting Reliability Goals
 - Elimination of Toil with Scripts and Automation
 - Monitoring Everything That Matters



3

Embracing Risk and Setting Reliability Goals



4

Reliability Goals

- This is the most **philosophical break** SRE makes **from traditional operations**, which often **aims** for **100% uptime**.
- **SRE acknowledges** that **100% availability** is **impossible** — it is **too expensive, slows innovation**, and **often unnecessary**.
- **Reliability** is treated as an explicit, **measurable goal**, not a vague wish.
- SRE manages the **trade-off** between **Agility** (*speed of releases*) and **Reliability** (*stability*) using **SLOs (Service Level Objectives)** and **Error Budgets** to balance reliability and agility.
- Reliability in SRE is built on four pillars:
 - **SLI** (What we measure)
 - **SLO** (What we target - goa)
 - **SLA** (What we promise customers - legal/commercial)
 - **Error Budget** (How much failure is allowed)



5



6

Reliability Goals | Service Level

- A Service Level refers to the specific, measurable **performance standard** or target that is agreed upon for the **delivery** of a **service**.
- It essentially answers the question: "**How well and how quickly will this service perform?**"
- **The Interconnected Relationship**
 - Think of them as a **hierarchy of commitment**:
 - **SLI: THE MEASUREMENT** (*The actual percentage of uptime, e.g., 99.95%*)
 - **SLO: THE TARGET** (*We aim for 99.9% Uptime*)
 - **SLA: THE COMMITMENT** (*If we fail to meet the 99.9% target, we pay you a penalty*)



7

Service Level | SLI

- **Service Level Indicators (SLIs)**
 - SLI is a **numerical measure** that tells **how reliable** the **system** is. The **raw data points** that tell you **what** the service is doing.
 - A **quantitative measure** of a **service's performance** (e.g., *99.99% of requests returned in under 100ms*). This detects risk.
 - SLIs **measure customer-impacting metrics**, such as:

SLI Type	What It Measures	Example Metrics
Availability	Service uptime, success response rate	% of successful API calls (2xx)
Latency	Speed of response	% of requests completed under 200 ms
Error Rate	Failed or erroneous transactions	5xx errors, failed login attempts
Throughput / Traffic	Requests handled per second	API requests/second, trades/sec
Saturation / Capacity	Resource exhaustion levels	CPU, memory, queue usage
Durability	Data integrity & loss protection	% of data loss-free operations
User Experience (UX)	Real user performance	App crash rate, page load time
Business SLI	Business/user success	% successful trades, payment success rate



8

Service Level | SLOs

- **Service Level Objectives (SLOs)**

- The agreed-upon **target reliability percentage** (e.g., *99.95% uptime*), based on SLI data.
- The explicit, **internal goal** for the reliability of a service.
- The SLO is an **internal contract** between the SRE team and the Development (Product) team.
- **Example**
 - The payment API should be 99.95% available over 30 days.
 - 95% of trades should complete within 300ms.
- It tells **how reliable** the service **SHOULD be**, not to be perfect, but **good enough for customers** at a **reasonable cost**.
- The SLO is the most crucial terms, as it **directly drives engineering behavior** and defines the Error Budget.



9

Service Level | SLOs ... continue

- SLOs are categorized based on **what type of reliability** we are measuring.

SLO Type	What It Measures	Example Metric	Example SLO Statement	Used For
Availability SLO	% of time service is available	Uptime %, Error rate	Service will be available 99.9% of the time per month	Service reliability, uptime commitments
Latency / Performance SLO	Speed of response	Response time (ms), Page load time	95% of API calls should respond under 200ms	User experience, app responsiveness
Throughput SLO	Volume of successful requests per time	Requests/sec, Transactions per minute	System should handle 5000 orders per minute	Scalability, system capacity
Error Rate SLO	Frequency of failed requests	% failed requests	Less than 0.5% of API requests can fail in a week	Quality of service, reliability
Durability SLO	Protection of data from loss	Data loss %	Data durability of 99.99999999% (11 nines) per year	Data storage, backup, cloud services
Reliability SLO	Successful task completion	Successful job completion rate	99.9% of scheduled jobs should finish successfully	Background jobs, workflows



10

Service Level | SLOs ... continue

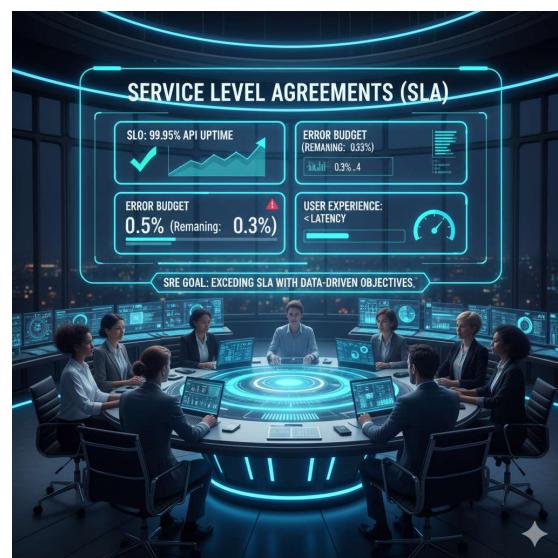
SLO Type	What It Measures	Example Metric	Example SLO Statement	Used For
Security SLO	Protection against threats	Vulnerability fix time	Critical security issues resolved within 72 hours	Security operations, compliance
Freshness / Data Staleness SLO	Timely data updates	Sync delay (mins), Data lag	Data replication lag must be under 5 minutes	Data pipelines, analytics
Capacity SLO	Handling load without degradation	CPU, RAM, storage usage	System must maintain < 70% CPU usage under peak load	Infrastructure scaling
Customer Experience SLO	Satisfaction & usability	CSAT, NPS, App rating	Maintain 90%+ customer satisfaction score	Business-level user experience



11

Service Level | SLA

- **Service Level Agreements (SLA)**
 - **Contract**
 - The **external, legal agreement** with the **customer**, outlining **penalties for failure to meet** the **SLO** (e.g., **99.5%** with a **10% credit** if missed).
 - **Buffer**
 - SLOs are typically set **tighter** than the **external SLA** to provide a **safety margin**, ensuring that if the team misses the **internal target** (the **SLO**), they still have a chance to fix it before breaching the **external legal contract** (the **SLA**).



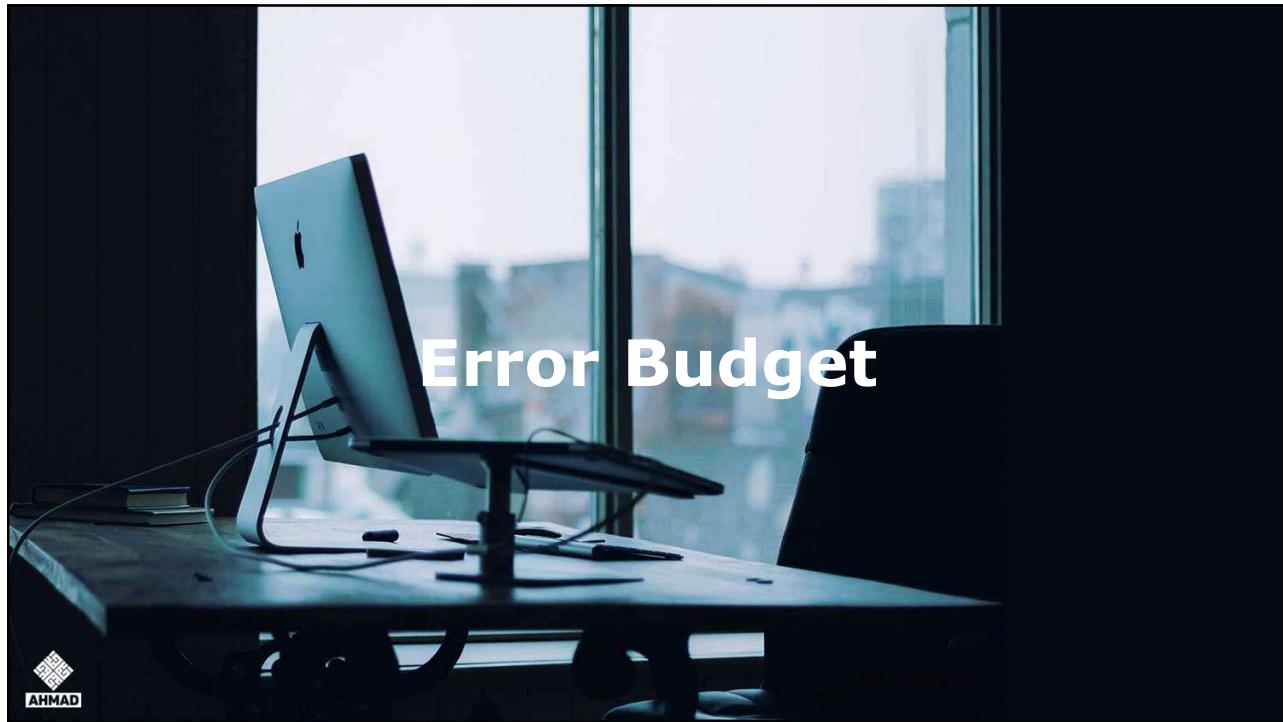
12

Service Level | SLA ... continue

SLA Type	What It Ensures	Real-Life Example
Availability SLA	Uptime guarantee	Server SLA: 99.99%
Support SLA	Response/resolution time	P1 issue responded in 15 min
Performance SLA	Response speed guarantee	95% of transactions < 200ms
Data Protection SLA	Backup, retention, durability	99.999999999% (11 nines) data durability
Compliance SLA	Security, regulation, audit	GDPR, PCI-DSS, RBI guidelines
Penalty-based SLA	Refund/credit if breached	Refunds 25% monthly bill



13



14

What is Error Budget

- The **Error Budget** is the **maximum amount of time a service is allowed to be unavailable** or fail over a **specific period** (e.g., one month) without violating its **Service Level Objective (SLO)**.
- It is fundamentally a **governance tool** used to **create alignment between Development teams** (who prioritize feature speed) and **SRE teams** (who prioritize reliability and stability).



15

Reliability Goals | Error Budget

- Error Budget**
 - The **acceptable amount** of time the **service is allowed to fail** (*calculated as 100% - SLO*).
 - Calculation**
 - The Error Budget is mathematically derived from the SLO and is **calculated** as the inverse of the SLO:
$$\text{Error Budget} = 100\% - \text{SLO}$$
 - Example:**
 - If your **SLO** (target reliability) is **99.9%** ("three nines"), your **Error Budget** is **0.1%**.
 - Over a **30-day period** (43,200 minutes), **0.1%** of the time **equals approximately 43 minutes** of acceptable **downtime** or error allowance.
 - Error budget formula:**

SLO (%)	Error Budget (%)	Max Allowed Failure Time (monthly)
99%	1%	~7.2 hours
99.9%	0.1%	~43.2 minutes
99.99%	0.01%	~4.3 minutes
99.999%	0.001%	~26 seconds



16

Reliability Goals | Error Budget ... continue

- o The Error Budget **provides** the SRE team with the necessary **authority** to manage risk:
 - **When the budget is full**
 - ✓ Dev teams can **move fast**, deploying new features frequently.
 - ✓ They are **embracing risk** because they have a **buffer**.
 - **When the budget is depleted (spent)**
 - ✓ All new **feature development** must **stop** (a feature freeze).
 - ✓ Engineering resources must be redirected to reliability work (*bug fixes, technical debt reduction*) until the budget is replenished.
- o **Error Budgets** can be **applied in different ways** depending on what is being measured:

Error Budget Type	Description	Example
Availability Budget	Allowed downtime	Max 43 minutes per month
Latency Budget	Allowed slow responses	No more than 1% calls > 500ms
Error Rate Budget	Allowed failed API requests	Up to 0.1% request failures
Data Loss Budget	Allowed inconsistency/drop	99.999% durability => error budget 0.001% loss
Performance Budget	Allowed performance drop	95% of orders processed under 3 seconds
Security Budget	Allowed exposure time / risk	Fix critical issues within 72 hours



17

Use Case | Stock Exchange

Concept	Meaning	Purpose	Stock Exchange Real-Life Example	Measurement/Criteria
SLI (Service Level Indicator)	Actual measured performance	Shows "how system is performing"	<ul style="list-style-type: none"> • 98.5% of trades completed within 50ms; • 0.02% error rate 	<ul style="list-style-type: none"> • Latency, uptime %, failure rate, data accuracy
SLO (Service Level Objective)	Target performance goal	Defines "how reliable system should be"	<ul style="list-style-type: none"> • 99.99% uptime during trading hours • <100ms trade processing time 	<ul style="list-style-type: none"> • Based on business demand & traffic
SLA (Service Level Agreement)	Legal/business contract with customer	Defines penalties if reliability is broken	<ul style="list-style-type: none"> • If uptime <99.95%, stock exchange owes financial compensation to brokers/banks 	<ul style="list-style-type: none"> • Uptime %, time-to-resolution • compensation terms
Error Budget	Allowed failure limit before reliability is affected	Balance reliability and innovation	<ul style="list-style-type: none"> • For 99.99% • SLO → allowed downtime = 4.3 min/month 	<ul style="list-style-type: none"> • Error Budget = 100% - SLO



18



19

Automation & Engineering

- The **Elimination of Toil** is a foundational principle of Site Reliability Engineering (SRE).
- It is the mechanism SREs use to move from a **reactive firefighting role** to a **proactive engineering role**.
- What is "Toil"**
 - Repetitive, manual, ticket-based, operational work** that adds **no long-term value**.
 - Examples of toil:**
 - Ticket processing
 - Log checking manually
 - Restarting failed services
 - Manual user provisioning

TOIL

SRE GOAL: AUTOMATING TOIL - 50% PROJECT WORK

20

Automation & Engineering ... continue

- SRE goal → **Automate repetitive work** to **free engineers** for **innovation**.
- **Characteristics of Toil**
 - **Manual** → Requires humans to perform it
 - **Repetitive** → Same process repeated frequently
 - **Reactive** → Triggered by alerts, tickets, or incidents
 - **No long-term value** → Does not improve the system
 - **Scale-dependent** → More workload, more manual effort
- **Key Practices:**
 - **The 50% Rule**
 - SREs must **spend no more than 50%** of their **time** on **toil**.
 - The other 50% is dedicated to Engineering (*writing code for automation, tools, and resilience*).
 - **Infrastructure as Code (IaC)**
 - Using tools like Terraform or Ansible to **manage infrastructure using code**, eliminating manual configuration.
 - **Automated Runbooks**
 - Turning manual incident **mitigation steps into scripts** that can be **triggered automatically**.



21

Types of Toil in IT Operations

Type of Toil	Description	Example in Industry
Monitoring Toil	Responding to repetitive alerts	Low CPU alerts at night
Deployment Toil	Manually deploying code/apps	Change requests executed manually
Configuration Toil	Manual server or firewall setup	SSH into each VM to update settings
Incident Toil	Handling same type of incidents repeatedly	Restarting crashed services daily
Reporting/ Documentation Toil	Creating manual weekly reports	Manually compiling uptime metrics
User/Access Management	Manually provisioning accounts	Creating SSH access requests for engineers



22

How Automation Eliminates Toil

Manual Operation (Toil)	Automated (Engineering-Based)
Manually restarting failed services	Self-healing automation using scripts
Logging into VMs to install patches	Automated patching via Ansible/Jenkins
Email-based change approvals	GitOps-based auto change control
Manual ticket creation after incident	Auto-ticket creation via ServiceNow API
On-call engineers checking alerts all night	Intelligent alert suppression and AI-based routing



23

Use Cases

- **Automated Payment Gateway Recovery (Bank)**
 - **Problem:** Payment system crashes frequently during **peak hours** → **manual restart** takes **30 minutes** → **revenue loss**.
 - **Automation Solution**
 - SRE creates an **auto-recovery script**:
 - ✓ Detects API failure
 - ✓ Restarts pods/servers automatically
 - ✓ Notifies engineers on Slack
 - **Manual time:** 30 mins → **Automated: 1 min**
 - **Result:** **80% incidents self-healed**, zero customer impact.
- **Automated Trade Order Scaling (Stock Exchange)**
 - **Problem:** During **market opening** (9:15 AM), **order load spikes 15x** → **website slow** → **manual scaling** takes **1 hour**.
 - **Automation Solution**
 - **Auto-scale** Kubernetes clusters using CPU, QPS, and latency triggers
 - Introduced **auto-deployment** verification **check**
 - **Manual scaling:** 1 hour → Auto-scaling: **2 minutes**
 - Zero revenue loss during market hours.



24

Tools Used for Toil Automation

Purpose	Tools Used
Infrastructure automation	Terraform, Ansible
Monitoring & Auto-healing	Prometheus, Grafana, AlertManager, PagerDuty
CI/CD automation	Jenkins, Spinnaker, GitHub Actions
Incident auto-resolution	Runbook Automation (StackStorm)
GitOps Change Management	ArgoCD, FluxCD



25

Monitoring Everything That Matters (Observability)



26



27

Monitoring

- Monitoring is the continuous process of collecting, analyzing, and visualizing data (metrics, logs, and traces) about a system to understand its behavior, detect problems, and inform decision-making.
- The modern approach to monitoring and system understanding, called Observability, relies on three primary types of data, often referred to as the **Three Pillars**:
 - **Metrics (What)**
 - **Logs (Why)**
 - **Traces (Where)**
- **Events** are a related, specialized type of **record** often captured **within logs** or treated as a distinct data source.
- Together, these data types provide a complete picture, answering the questions: **What is happening? Where is it happening? and Why is it happening?**



28

Metrics: The Quantitative View ("What")

- Metrics are **numerical measurements** that represent the **health** and **performance** of a system over time.
- They are the most efficient way to track trends and trigger alerts.
- Metrics are structured, **time-series data**: a **number collected** at a **specific timestamp** and enriched with descriptive labels (e.g., service name, region).
 - Granularity: Aggregated**
 - Metrics are **summarized values** (averages, sums, percentiles) over a **time interval** (e.g., 99th percentile latency over the last minute).
 - Purpose**
 - Ideal for **monitoring dashboards**, setting **Service Level Objectives** (SLOs), and generating **real-time alerts** when a threshold is breached (e.g., CPU usage is 80%).
 - Examples**
 - Response time/Latency, Request Rate, Error Rate, CPU/Memory Utilization.



29

Logs: The Detailed Narrative ("Why")

- Logs are **immutable, timestamped text records** of discrete events that occurred within an application or system. They are the system's diary, providing rich contextual detail.
- Textual data**, often **unstructured** or **semi-structured** (like JSON), **generated** when **code executes** or a **system action occurs**.
- Granularity: Granular and Contextual**
 - Each line typically refers to a **single event** with **specific details**, including the **exact state** of the **program** at that **moment** (e.g., variable values, stack traces).
- Purpose**
 - Essential for **root cause analysis**, **forensic debugging**, and **auditing**.
 - When a metric alerts, SREs turn to the logs to find the "why."
- Examples**
 - [ERROR] NullPointerException in payment service at line 45
 - [INFO] Database connection established
 - [WARN] Cache miss for user ID 12345.



30

Traces: The Request Journey ("Where")

- Traces, or **Distributed Tracing**, map the **entire life cycle** of a **single request** or **transaction** as it **moves** through a complex, distributed **system** (like microservices).
- **Span**
 - A span is a **single, logical unit of work performed** by a **service** or **component**, such as:
 - An incoming or outgoing **HTTP request**.
 - A **database query**.
 - A **specific function execution**.
- A **trace** is a **collection** of **spans**. All **spans belonging** to the **same request** share a unique **Trace ID**.
- **Granularity: End-to-end journey**
 - They link together the **operations across service boundaries**.
- **Purpose**
 - Crucial for **diagnosing latency and performance bottlenecks** in microservices.
 - It **visually shows** which **service** is **responsible** for most of the **delay**.
- **SRE Use**
 - When a latency metric spikes, the trace for a slow request immediately points to **which service** in the call graph introduced the delay.



31

Events: The State Change ("When")

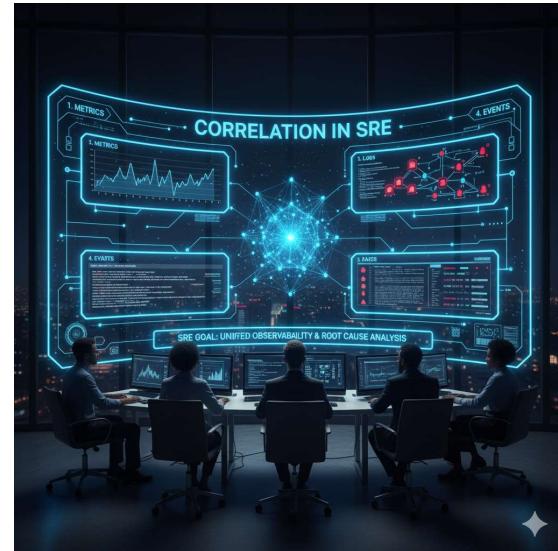
- While **often captured within logs** or used to derive metrics, **Events** are distinct occurrences that **represent** a significant, **discrete action** or **change of state** in a system.
- Structured records of key incidents or changes, **generated** by both **systems** and **humans**.
- They are generally of higher importance than routine log messages.
- **Granularity: Discrete Occurrence**
 - They **mark a specific moment in time**.
- **Purpose**
 - Primarily used for **correlation** and **context**.
 - They **answer** the question: "**What external action occurred right before the problem started?**"
- **Examples**
 - DeploymentStarted (version 3.1.2)
 - ServiceScaledUp (3 to 5 instances)
 - ConfigurationChanged (feature flag X disabled)
 - SecurityAlertTriggered



32

The Value of Correlation

- The true power of Observability comes from **correlating** these data types.
- **For instance:**
 1. A Metric (**Request Rate**) drops suddenly.
 2. The SRE checks **Events** and sees that a **DeploymentStarted** event occurred seconds before the drop.
 3. The SRE then pivots to **Traces** for the **affected service**, seeing a **high failure rate**.
 4. Finally, the SRE drills into the **Logs** linked to the failed traces, finding the exact **stack trace** from the new code that caused the issue.
- This workflow **dramatically speeds** up the **Mean Time To Resolution** (MTTR).



33

Defining Meaningful Metrics and Alerts

- **Effective monitoring** starts with choosing **what to measure** and **deciding** when to **raise an alarm**.
- The goal is to maximize **signal** (*real, actionable problems*) and minimize **noise** (*irrelevant alerts*).
- **Meaningful Metrics**
 - Metrics are numerical, time-series data points that **quantify system performance**.
 - **Meaningful metrics align** with your **business goals** and directly **reflect** the **user experience**.
- **Effective Alerts**
 - An alert is a notification that a metric has crossed a predefined threshold or pattern, indicating a problem that requires attention.



34



35

White-Box vs Black-Box Monitoring

- These **two concepts** describe the perspective from which you **monitor a system**. A **complete monitoring strategy** typically **combines both**.
- **Black-Box Monitoring**
 - Black-box monitoring **treats** the **system** as an **opaque**, closed-off entity—a "**black box**".
 - It is symptom-oriented and **checks externally visible behavior**.
 - **Perspective**
 - ✓ The **end-user's point of view**.
 - **Data Source**
 - ✓ **External probes**, synthetic **transactions** (*simulated user requests*), and **external system calls** (e.g., *ping*, *HTTP checks*).
 - **Focus**
 - ✓ **What is broken?** (*Is the service available, fast, and correct?*)
 - **Example Metrics**
 - ✓ **Uptime/Availability** (e.g., *is the login page returning a \$200\$ status code?*).
 - ✓ **External Latency** (*how long a complete user transaction takes from outside the network*).
 - ✓ **DNS Resolution** Time from a global location.
 - **Purpose**
 - ✓ To confirm the **Service Level Objective** (SLO) — the promised level of service — is being met.
 - ✓ It ensures that **regardless of internal health**, the **system is working for the customer**.



36

White-Box vs Black-Box ... continue

- **White-Box Monitoring**
 - White-box monitoring **uses information exposed** by the **internals** of the **system** — it **looks inside** the "white box".
 - It is **cause-oriented** and requires **direct instrumentation** of the **application code** or **host operating system**.
 - **Perspective**
 - The **system's internal** components and code.
 - **Data Source**
 - **Application logs, custom metrics** exposed by the **code (instrumentation)**, and **operating system/agent** data.
 - **Focus**
 - **Why is it broken?** (*What is the root cause: high memory, database lock, slow function call?*).



37

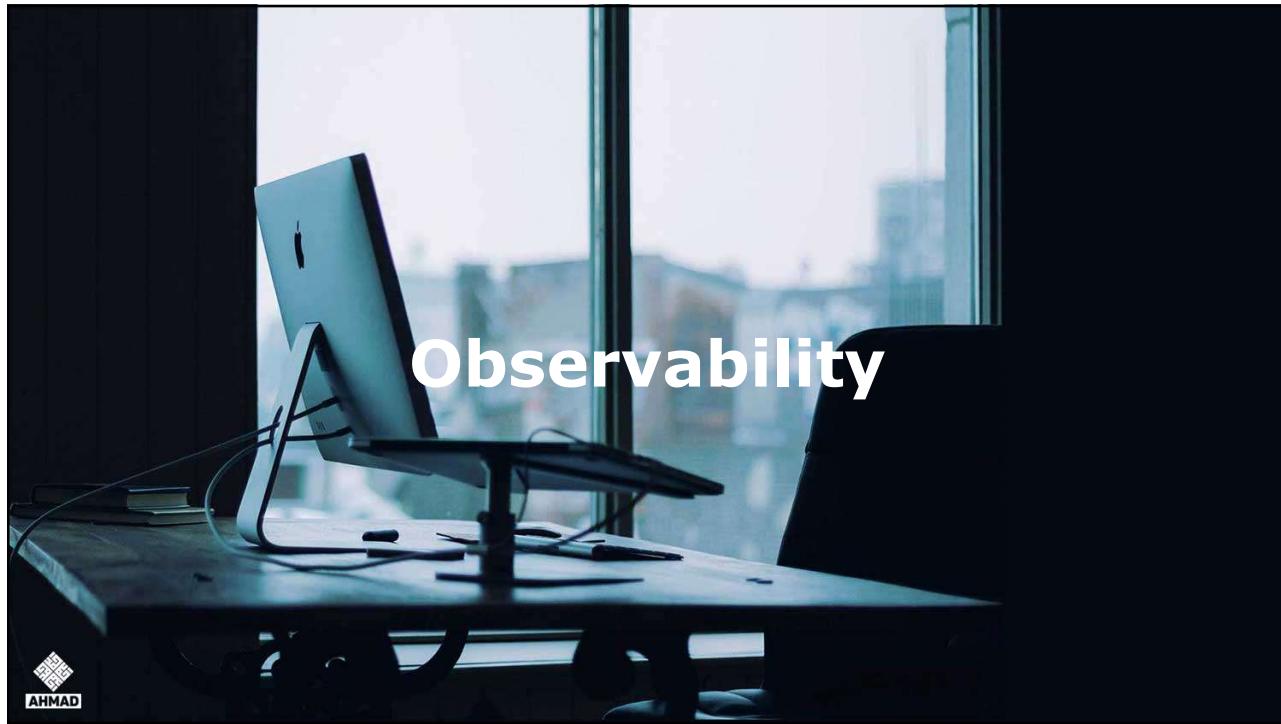
White-Box vs Black-Box ... continue

- **Example Metrics**
 - Internal CPU/Memory/Disk Usage of a specific server.
 - Garbage Collection frequency and duration.
 - Internal Queue Lengths or Thread Pool Utilization.
 - Specific function execution times measured within the application code.
- **Purpose**
 - To provide the granular detail needed for **debugging, capacity planning, and proactive issue detection** before a complete user-visible failure occurs.

Aspect	Black-Box Monitoring	White-Box Monitoring
View	External (User's Perspective)	Internal (System's Perspective)
Focus	Symptoms (Availability, Latency)	Causes (Resource Utilization, Code Logic)
Action	Reactive (Service is currently down/slow)	Proactive & Debugging (Service may go down; find the root cause)



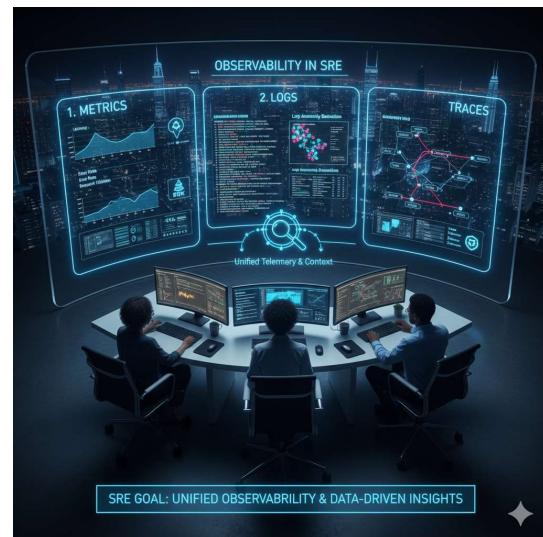
38



39

Observability

- SRE recognizes that traditional monitoring (*checking if a server is alive*) is insufficient for complex distributed systems.
- You must monitor what the user *experiences*.
- The shift from generic monitoring to **Observability**, which means collecting and analyzing data that allows engineers to ask novel questions about the system without knowing the answer beforehand.
- Unlike **traditional monitoring**, which answers "*Is the system up?*", **Observability answers** "*Why is the system slow right now, and which component is causing it?*"



40

Observability | Principles

- SREs prioritize **monitoring** the system based on the **user experience** and **business impact**.
- The **Four Golden Signals** are the definitive **set of metrics** that **must** be **tracked** for every user-facing service:

Signal	Definition	Why it Matters (User Impact)
Latency	The time it takes to serve a request .	High latency means a slow user experience, frustrating customers .
Traffic	A measure of how much demand is being placed on your service (e.g., requests per second).	Used for capacity planning and determining if a latency spike is due to high load .
Errors	The rate of requests that fail (explicitly via 5xx errors or implicitly via bad data).	Direct measure of user-facing failure. Every error spends the Error Budget .
Saturation	How full the service is; a measure of system utilization (<i>CPU, memory, disk I/O, queue depth</i>).	Predicts future failure. High saturation means the system is about to degrade performance or crash.



41

Observability | Pillars

Pillar	Purpose	What It Helps Identify	Tools & Tech	Real Use Case (Industry)	Benefits
Metrics	Collect numeric performance indicators over time	Resource usage , performance trends, saturation (capacity), success / failure ratios	Prometheus, CloudWatch, Datadog, Grafana	Stock Exchange: Monitor matching engine latency & CPU usage to ensure trade orders execute under 50ms	Proactive detection, trend analysis, capacity planning
Logs	Record detailed events, system messages, and error details	Root cause of failures , system bugs, security attempts, transaction history	ELK (Elastic, Logstash, Kibana), Splunk, Fluentd	Banking: Track failed payment logs, fraud detection logs, ATM withdrawal issues	Helps in debugging, audit trails, compliance, RCA
Traces	Follow a request across microservices (end-to-end tracking)	Identify bottlenecks , slow microservices, API dependency issues	OpenTelemetry, Jaeger, Zipkin	E-Commerce: Track order from cart → payment → inventory → shipping service and find where delay happens	Deep visibility into distributed systems
Events / Alerts	Notify when thresholds or anomalies occur	Unexpected failures, security breaches, system outages	PagerDuty, Prometheus Alert Manager, Opsgenie,	Healthcare System: Alert if medical record system response time exceeds 2 seconds	Real-time detection & response



42

Observability | Pillars ... continue

Pillar	Purpose	What It Helps Identify	Tools & Tech	Real Use Case (Industry)	Benefits
Business KPIs (Service Health)	Track business impact instead of just system health	Customer loss, transaction drop , revenue impact, churn rate	Datadog, New Relic, Google Analytics, Dynatrace	Stock Exchange: Number of failed trade settlements or delayed transaction confirmations	Bridges business & technical metrics, aligns SRE with business goals
User Experience Monitoring (RUM & Synthetic)	Measure real user interactions & experience	Page load time , mobile app crash, sign-in failure	New Relic, Dynatrace, Datadog RUM, Pingdom	Online Banking: Detect when customers struggle with login or fund transfer due to UI lag	Captures true user experience; helps identify frontend issues
Security Monitoring	Monitor threats, suspicious access, vulnerabilities	Unauthorized access , DDoS, fraud detection	SIEM, Splunk Enterprise Security	Banking: Detect brute force login attempts or fraudulent ATM transactions	Protects reliability, compliance, trust
Infrastructure Monitoring	Monitor servers, network, storage	VM crash, disk full, network latency	Nagios, Zabbix, Datadog Infra	Telecom Industry: Monitor 5G tower network connectivity, loss, signal strength	Ensures platform stability and environment reliability



43

Observability | Benefits

Benefit	Explanation
Faster Incident Detection	Reduces MTTR (Mean Time to Recovery)
Predictive Reliability	Detect failures before they impact customers
Prevents SLA Breaches	Helps stay within error budget
Enables Auto-Remediation	Integrates with tools like StackStorm
Helps SRE, DevOps, and Business	Unifies technical & business insights



44



45

What is Full-Stack Observability

- Full Stack Observability provides **deep, holistic visibility** into the performance and behavior of an application, from the **end-user interaction** all the way down to the **underlying network and infrastructure**.
- It integrates the three pillars of observability: **Metrics, Logs, and Traces**, across every component of the system.



46

Beyond Infrastructure

- **Full Stack Observability** ensures **comprehensive monitoring** across three critical layers that are often siloed in **traditional monitoring** setups:
 - **Application (App) Visibility**
 - Focuses on the **code** and **services** that make up the **application** (e.g., *microservices, APIs, databases*).
 - **Key Data: Traces, Metrics and Logs.**
 - **Network Visibility**
 - Focuses on the **data flow between application components**, users, and the cloud/data center.
 - **Key Data:**
 - **Network latency, packet loss, bandwidth utilization, firewall performance, and load balancer health.**
 - This is crucial for modern distributed systems where network slowness can often be mistaken for application issues.



47

Beyond Infrastructure

- **User Experience (UX) Visibility**
 - Focuses on the actual experience of the end-user, often through **Real User Monitoring (RUM)** and **Synthetic Monitoring**.
 - **Key Data:**
 - **Frontend Performance**
 - **Page load times**, core web vitals (*Largest Contentful Paint, Cumulative Layout Shift*), JavaScript **errors**, and **interaction latency**.
 - **Business Impact**
 - **Conversion rates**, cart abandonment rates, and **other key business metrics** tied to performance.



48

Data for Root Cause Analysis (RCA)

- The true power of **Full Stack Observability** lies in its ability to **correlate data across** the entire **stack**.
- Enables Better RCA:**
 - Unified Data Platform**
 - All **Metrics**, **Logs**, and **Traces** are ingested and linked within a **single platform**.
 - Contextual Tracing**
 - A **single user request** is **tracked** (via a Trace ID) from the user's browser, through the load balancer, across multiple microservices, and **down** to the **database query**.
 - If a database query is slow, the trace immediately highlights the specific service and the offending query, linking it to associated error logs.
 - Automatic Correlation**
 - The platform **automatically identifies dependencies** and **anomalies**.
 - For example, if **Application Metrics** show a **spike in errors**, the FSO system can automatically **check if Network Metrics** show **increased latency** or if Infrastructure Metrics show a **spike in CPU utilization** at the exact same time.
 - Faster MTTI and MTTR**
 - This correlation dramatically **reduces** the **Mean Time To Identify** (MTTI) the component causing an incident and the **Mean Time To Resolve** (MTTR) it, as SREs spend less time navigating tool silos.



49



50

Balancing Velocity with Reliability

- The core purpose of an **Error Budget** is to serve as a **risk regulator**, helping **balance** the ongoing tension **between** moving **quickly** (*agility and velocity*) and maintaining **stability** (*reliability*).
- Feature Velocity**
 - Developers are incentivized to **release features quickly** and often, which inherently introduces risk.
- Budget as a Buffer**
 - The Error Budget acts as a buffer.
 - As long as the **budget is healthy** (full), the teams can maintain **high velocity**, knowing they have room to absorb minor bugs or failures.
- This mechanism ensures that reliability commitments are **enforced by data** rather than endless meetings or subjective arguments.



51

Drive Engineering Decision

- The Error Budget is a powerful tool for strategic prioritization:
 - Prioritizing Toil Elimination**
 - SREs can **justify spending time** on **automation** by calculating the potential **budget savings**.
 - Example**
 - If **manual configuration** causes one **10-minute outage per month** (*spending 10 minutes of the budget*), **automating it permanently justifies** taking one week off feature work to write the automation code.
 - Guiding Test Coverage**
 - If a **critical service** has a **very tight SLO** (e.g., 99.999%), its **Error Budget** is **minuscule**.
 - This drives an engineering decision: **Invest heavily in automated testing** and staging environments to **ensure changes** are virtually **risk-free before deployment**.



52

Drive Engineering Decision ... continue

- **Informing System Design**
 - Services with **large, unused budgets** are **prime candidates for de-scoping** (*reducing over-engineering*).
 - Services that **constantly burn their budget** need urgent investment in **redesign** or architectural changes to **improve resilience**.
- **Blameless Learning**
 - After an **incident**, the post-mortem **focuses** on **how much budget was spent** and **why**.
 - The key action item is always an engineering task to **prevent the budget-spending** incident from recurring, thus driving a culture of continuous improvement.



53

Help in Decisions

Scenario	Decision
Budget mostly unused	Approve new releases, experiments, A/B testing
Budget moderately used	Approve with caution, implement canary deployment
Budget nearly exhausted	Change freeze, focus on performance & bug fixes
Budget completely consumed	Block deployments, trigger high-priority reliability review



54



SRE Roles & Responsibilities

55

SRE Roles and Responsibilities

- An SRE is a software engineer who possesses deep operational knowledge.
- Their time is rigidly split between managing the current system and engineering future improvements.

Responsibility	Traditional IT Ops Approach	SRE Approach
Monitoring	Server uptime, CPU, memory	Service health, SLO, SLI, latency, error rate, business impact
Incident Management	Reactive resolution	Proactive prevention , blameless postmortem, error budget-based decisions
Change Management	Manual approvals, CAB meetings	Automated CI/CD-based GitOps, canary, blue/green deployments
Troubleshooting	Manual logging, ticket-based	Automated root-cause analysis, observability, distributed tracing
Capacity Planning	Static provisioning	Auto-scaling , performance benchmarking, chaos engineering
Toil Handling	Manual repetitive tasks	Automation -first approach (Terraform, Ansible, scripts, AI Ops)
Reliability Engineering	NA (mostly infrastructure focus)	Error budgets, SLO/SLA , reliability testing, chaos engineering



56

SRE Team

- An **SRE (Site Reliability Engineering) team** is a specialized group of engineers who focus on **making systems reliable, scalable, efficient, and resilient** — using **software engineering, automation, observability, and proactive problem-solving** rather than manual operations.
- Types of SRE Engineers:**

SRE Role	What Work They Perform (Responsibilities & Tasks)	Tools & Technologies
Infrastructure SRE	Manages cloud and hybrid infrastructure , Kubernetes, IaC, network, provisioning, reliability engineering, cost optimization	AWS, Azure, GCP, Kubernetes, Terraform, Ansible, Helm, VMware, OpenStack, Ansible, Bash/PowerShell, Nagios, SolarWinds, Zabbix, NetApp
Platform SRE	Builds internal developer platforms , CI/CD, Observability, Secrets, Service Catalog, GitOps, SSO, API gateway	GitHub Actions, Jenkins, ArgoCD, Vault, ServiceNow, Grafana, ELK
Product/Application SRE	Ensures reliability of user-facing apps (payment, mobile banking, e-commerce, trading), sets SLO/SLIs, monitors production health, chaos testing, RCA	Datadog, Splunk, Jaeger, PagerDuty, New Relic, Chaos Monkey
Security SRE (Sec-SRE)	Threat monitoring, vulnerability management , automated patching, compliance, IAM, DDoS protection, security event automation	Splunk SIEM, Palo Alto, Cloudflare, Qualys, AWS GuardDuty
ML/AI SRE (MLOps)	Monitor reliability of ML models , automate retraining, manage data pipelines, ensure model accuracy & performance	Kubeflow, MLflow, TensorFlow, SageMaker, Kafka
Database SRE	DB availability , replication, backup, auto-recovery, sharding, query tuning, HA clusters, DR testing	Oracle RAC, PostgreSQL, Cassandra, MySQL, MongoDB, AWS RDS



57

How SRE is Structured in Organizations

- 3 Major SRE Structures:**

SRE Structure Type	Description	Company Examples	When to Use
Embedded SRE	SRE engineers directly work inside product teams	Google Search, Amazon Prime	Complex large-scale products needing continuous reliability
Central SRE / Shared SRE Team	Single SRE team supports multiple applications	Banking, telecom, SaaS	Medium-scale apps, central governance
Platform SRE	Works on common tools , CI/CD, infrastructure, observability, SLO frameworks	Netflix, Spotify, Salesforce	When building developer experience or DevOps platform



58

What Skills Do SRE Engineers Need

Skill Category	What SRE Should Know	Example
Software Development	Python, Go, Shell scripting, Git	Build automation, monitoring scripts
Cloud & Infra	AWS, Azure, GCP, Kubernetes, Terraform	Deploy Google Kubernetes, scale AWS infra
Automation & DevOps	CI/CD (Jenkins, GitHub), Ansible, GitOps, IaC	Automated deployment
Observability & Monitoring	Prometheus, Grafana, ELK, Splunk, Datadog	Track latency, error rates, business KPIs
Reliability Concepts	SLO, SLI, SLA, Error Budget, Blameless Postmortem	Decide when to stop deploying
Chaos Engineering	Gremlin, Chaos Monkey	Test failure recovery
Incident & Problem Solving	Root Cause Analysis, self-healing	Analyze high CPU issue, create fix
Performance & Scalability	Load testing, auto-scaling	Ensure 1M users during sale
Soft Skills	Collaboration, communication, problem-solving	Talk to developers and business



59

Roadmap to Convert Traditional IT Ops to SRE

Phase	Goal	Activities
Assessment	Identify skill gaps	Evaluate current Ops skills, mindset, tooling
Training	Build new capabilities	Train on Python, Cloud, CI/CD, Monitoring, SRE fundamentals
Automation First	Reduce toil	Transition from manual tickets to automation & scripts
Introduce SLO / SLA / Error Budgets	Reliability focus	Implement reliability metrics, observability dashboards
Create SRE Squads	Start hybrid model	Mix Dev + Ops + SRE together (Platform, Infra, Product SRE)
Full SRE Adoption	Establish culture	Blameless postmortems, continuous improvement, chaos testing



60

Ops to SRE Transformation

Before SRE (Ops Team)	After SRE (Reliability Engineering Team)
Manual monitoring of servers	Automated trade monitoring using latency, TPS, failure rate
Ticket raised for payment failure	Script auto-restarts payment service (self-healing)
Outage analysis done manually	SRE uses distributed tracing (Jaeger) to find bottleneck
Change approval in email	GitOps-based approval and automated deployment
Manual scaling during "Salary Day"	Auto-scale compute based on trade volume



61

Case Study



62

Employee Identity Service (EIS)

- Your organization (**FinTech Corp**) has an **Employee Identity Service (EIS)** responsible for:
 - Employee **authentication** (login)
 - **Authorization** (role-based access control)
 - **Profile retrieval** (name, designation, department)
- This service is **mission-critical**, powering:
 - **Payroll & HR Portal** → Salary slips, leave & appraisal
 - **Finance Approval System** → Expense approvals
 - **IT Service Desk** → Asset provisioning & access requests
 - **Onboarding Platform** → New employee setup
- **Business Expectation**
 - Employees should be able to access portals securely, quickly, and reliably without login issues, especially during salary processing, appraisal cycles, and onboarding.



63

Step 1: Define Possible SLIs

SLI Type	Example Metric	Description
Availability		
Latency		
Correctness		
Error Rate		
Freshness		



64

Possible SLIs

SLI Type	Example Metric	Description
Availability	% of successful login requests	Measures uptime/access success
Latency	% login responses < 500ms	Measures login response speed
Correctness	% correct identity responses	Ensures data accuracy (role, name)
Error Rate	% failed login attempts (5xx)	Tracks failures due to service issues
Freshness	Time since last user update synced	Measures how current data is



65

Step 2: Define Possible SLOs

SLI	SLO (Target)	Error Budget
Availability		
Authentication Latency		
Error Rate (5xx)		
Correctness		
Freshness		



66

Possible SLOs

SLI	SLO (Target)	Error Budget
Availability	99.5% monthly	0.5% downtime = 216 mins/month
Authentication Latency	95% requests < 400ms	5% slow requests allowed
Error Rate (5xx)	<0.3% per month	0.3% failure margin
Correctness	99.9% accurate identity responses	0.1% incorrect allowed
Freshness	95% of syncs within 1 hour	5% can exceed 1hr



67

Step 3: Define Possible SLA

SLA Parameter	SLA Commitment
Availability	
Login Response Time	
Identity Data Sync	
Incident Response	
Compensation	



68

Possible SLA

SLA Parameter	SLA Commitment
Availability	99.0% quarterly (<i>max downtime: 1h 48m per month</i>)
Login Response Time	90% of login responses < 500ms
Identity Data Sync	Max sync delay = 4 hours
Incident Response	Critical incidents resolved within 4 hours
Compensation	Service credits: e.g., 5% monthly subscription refund for downtime exceeding SLA



69

Discussion

- Why SLO is stricter than SLA
- What happens if SLA is breached externally
- Error Budget vs SLA



70

Discussion | Answer

- **Why SLO is stricter than SLA**
 - Internal target (SLO) is higher to **ensure service reliability**
 - SLA is slightly lower to **protect contractual commitments** and allow operational tolerance
- **What happens if SLA is breached externally**
 - Financial **service credits** (refund, subscription adjustment)
 - **Priority support / escalation** for affected customers
 - Reputation & trust may be impacted → triggers postmortem and reliability improvements
- **Error Budget vs SLA**
 - **SLA: 99.9%** uptime
 - **Error budget: 0.1% allowed downtime** per month = 43 minutes of downtime allowed per month



71



72

National Stock Exchange

Traditional Ops Environment	Modern SRE Transformation
Manual monitoring via NOC	Full observability (<i>Grafana, Prometheus, AI alerting</i>)
1–2 hours downtime during trading time	Self-healing failover in 2 minutes
1000+ monthly tickets	80% tickets eliminated via automation
Change freeze during market hours	Safe, automated deployments using canary
Problem Resolution = War-room	Resolution via Runbooks , automated fixes



73

Why Change Was Needed

- One **trading outage** caused a **₹100 Crore loss** and **reputation damage**.
- Analysts were **fixing** production **incidents manually**.
- **Deployments** were **done via scripts**, after midnight, with **risk and fear**.
- **Monitoring** was **limited** to **CPU & RAM** — no business metrics.



74

How They Introduced SRE

Phase	Step	SRE Activities
Phase 1	Analyze	<ul style="list-style-type: none"> Define SLOs (e.g., 99.99% uptime) Identify most frequent failures
Phase 2	Automate	<ul style="list-style-type: none"> Incident detection → automated alerts Self-healing scripts → restart crashed services
Phase 3	Reduce Toil	<ul style="list-style-type: none"> 70% tickets converted to automated jobs Password resets, patching, backup monitoring
Phase 4	Build Reliability Culture	<ul style="list-style-type: none"> Introduced blameless postmortems Used error budget policy for release decisions
Phase 5	Optimize	<ul style="list-style-type: none"> Chaos testing, load simulators, active failover



75

Results of the SRE Transformation

Before SRE	After SRE
4 major outages/year	0 major outages in 12 months
Recovery Time: 1–2 hours	Self-healing in under 3 minutes
1000+ tickets/month	80% reduction (automated resolution)
6–8 people required for manual patching	Fully automated patch pipeline
No visibility into root cause	Blameless postmortems & RCA automation



76

Final Summary

Traditional Ops	SRE Mindset
Block changes to maintain stability	Engineer systems to handle change safely
Reactive firefighting	Proactive & predictive reliability
Manual troubleshooting	Automation & self-healing
Ticket-based operations	Engineering-based operations
Support function	Strategic reliability engineering



77

Q & A

Any concepts still unclear?

Thank you for attending



78

39