

# Platform Engineering, AIOps & Continuous Improvement

1


## Platform Engineering

2

1

# What is Platform Engineering

- Platform Engineering is the practice of **designing, building, and maintaining internal platforms** that **provide engineers** with self-service **capabilities** — such as **automated provisioning, CI/CD workflows, observability, security, and compliance**— through standardized **tools, templates, and processes**.
- It bridges the gap between development and operations by providing **reliable, consistent, and automated platforms** that accelerate delivery while ensuring reliability.



3

## Platform Engineering ... continue

- It is characterized by:
  - **Product Mindset**
    - The platform itself is treated as an **internal product**, with **developers** as its **customers**.
    - The platform team **gathers feedback, manages a roadmap**, and aims to deliver a cohesive, user-friendly experience.
  - **Abstraction and Self-Service**
    - The platform team **provides managed services** and **APIs** that allow **developers** to provision infrastructure, deploy code, and set up monitoring.
  - **Standardization and Consistency**
    - By creating **standardized tools** and **workflows**, the platform ensures that all applications are built, deployed, and run in a consistent, compliant, and secure manner across the organization.



4

# Platform Engineering Supports SRE

- Platform Engineering enables SRE by **embedding reliability practices, automation, and governance into standardized internal platforms** that development and operations teams use.
- **Platform Teams as Enablers of Reliability and Consistency**
  - Platform teams **do not directly own application** uptime.
  - Instead, they **enable reliability** by:
    - **Standardizing operational practices** – logging, monitoring, deployment, rollback, alerting, backup, cost tracking.
    - **Embedding SRE principles** into platform capabilities such as auto-scaling, fault tolerance, resilience testing, and error budget policies.
    - **Ensuring consistency** across environments (Dev, QA, Prod) with reusable templates, Infrastructure as Code (Terraform, CloudFormation, Pulumi).
    - **Reducing TOIL** by automating repetitive tasks: provisioning, compliance checks, configuration, ticket routing.
  - **In SRE context:**
    - **Platform engineering ensures** that **every application inherits reliability best-practices by default** — without manually reinventing them for each service.



5

# Platform Engineering Supports SRE ... continue

- **Designing Reusable Automation Frameworks and Service Templates**
  - Platform engineers **create reusable components** called **Golden Paths, Service Templates, or Paved Roads**, which developers can use to deploy new services with built-in reliability.
  - These **templates** typically **include**:

Component	Description	SRE Benefit
<b>Terraform/Helm Templates</b>	Preconfigured compute, storage, networking, Kubernetes resources	Consistency, reliability, faster onboarding
<b>CI/CD Pipelines</b>	Automated build, test, deployment with quality gates	Reduces deployment risk, ensures change reliability
<b>Monitoring Blueprint</b>	Predefined SLI/SLO dashboards, alert rules, tracing setup	Built-in observability from Day 1
<b>Security Governance</b>	IAM, encryption, policy compliance checks	Reduces vulnerabilities and operational risk
<b>Incident Automation</b>	Auto-rollback, alert routing, runbook triggers	Faster recovery, lower MTTR




6

## Platform Engineering Supports SRE ... continue

- **Example:**
  - A Platform team provides a “**Service Deployment Template**”.
  - Developers select a **template** and automatically get:
    - SLO monitoring (latency, error rate)
    - Canary deployment
    - Auto-restart and self-healing
    - Logging and tracing integration
    - Error budget tracking
- No manual setup — *reliability is built-in*.
- Platform Engineering **empowers SRE practices by design**, making reliability, consistency, and automation *default rather than optional*.
- It transforms operations from “manual firefighting” to **reusable engineering**.



7

A dark blue background featuring a glowing, circular pattern of white and yellow particles, resembling a celestial body or a microscopic view of data. The word "AIOps" is centered in the middle of this glowing ring.

AIOps



8

# What is AIOps

- AIOps (**Artificial Intelligence for IT Operations**) is a multi-layered technology platform that combines **Big Data** and **Machine Learning (ML)** to automatically analyze the massive, complex data streams generated by IT infrastructure, applications, and performance monitoring tools.
- Its primary goal is to **automate and enhance IT operations** by turning reactive, manual processes into proactive, intelligent, and autonomous workflows.
- **Tools:** BigPanda, Moogsoft, IBM Watson AIOps



9

# AIOps | Types

- AIOps solutions are generally **categorized** by the **breadth** of their **focus**:
  - **Domain-Centric AIOps**
    - **Focused** on a specific, **single IT domain** (e.g., *Network Performance Monitoring*, *Application Performance Monitoring (APM)*, or *Security*).
    - These tools use **AI/ML tailored** to a narrow set of **data types** and **problems within** that **domain**.
  - **Domain-Agnostic AIOps**
    - Designed to **integrate data** from across the **entire IT stack** (*logs, metrics, events, topology, change data, etc.*) regardless of the source.
    - These platforms **correlate alerts** and **find root causes** across **network, application**, and infrastructure boundaries.



10

# AIOps | Capabilities

AIOps	Description	Benefit	Example in Real Scenario
Anomaly Detection	Detects <b>unusual patterns, sudden metric spikes, or abnormal system behavior</b>	<b>Prevents silent failures</b> and unexpected outages	AI detects <b>unusual memory consumption</b> in a database server <b>at 2 AM (much higher than usual)</b>
Event Correlation	<b>Groups related alerts</b> from multiple tools into a single meaningful incident	<b>Reduces alert fatigue</b> and noise	<b>150 alerts from servers, DB, and network</b> are automatically correlated into one major <b>incident</b> : " <i>Payment service API failure caused by DB connection timeout.</i> "
Root Cause Analysis (RCA)	<b>Identifies the most probable cause</b> of failure using historical and dependency data	Faster troubleshooting and <b>MTTR reduction</b>	AI <b>analyzes event history</b> and reports: " <i>Likely root cause: recent change in firewall policy blocking API traffic.</i> "
Predictive Analytics	<b>Uses ML to forecast failures</b> , performance degradation, or capacity exhaustion	Enables <b>proactive incident prevention</b>	AI <b>predicts</b> that <b>storage will reach 90% capacity in 12 days — triggers auto-scale</b> and sends proactive alert
Auto-Remediation	<b>Automatically resolves recurring or known issues</b> using scripts or workflows	<b>Reduces manual toil</b> and speeds up resolution	When a service becomes unresponsive, AI <b>automatically restarts</b> the Kubernetes pod or clears cache without human intervention
Business Impact Analysis	<b>Maps incidents</b> to business services, <b>SLAs</b> , or revenue impact	Helps <b>prioritize critical incidents</b>	AI <b>detects</b> that <b>latency issues in login API</b> are impacting 8,000 users and <b>may breach SLA</b> — marks it as <i>high priority business incident</i>



11

# AIOps | Benefits

- The core value of AIOps lies in its ability to manage the scale and complexity of modern, cloud-native environments that overwhelm human operators.
- Alert Noise Reduction:** ML models correlate thousands of raw, noisy events and alerts into a few **actionable incidents**. This drastically reduces **alert fatigue** for on-call engineers.
- Faster Root Cause Analysis (RCA):** AIOps quickly analyzes system dependencies and change data to pinpoint the most likely root cause, significantly lowering the **Mean Time To Resolution (MTTR)**.
- Proactive/Predictive Detection:** By learning normal behavior patterns, AIOps can perform **anomaly detection**, flagging subtle deviations *before* a traditional static threshold is breached, preventing outages rather than reacting to them.
- Automation:** Automates repetitive tasks, from triaging and routing tickets to executing self-healing runbooks (e.g., restarting a service or scaling a resource).
- Reduced Operational Costs:** Increased automation and faster incident resolution reduce the time SREs and Operations teams spend on firefighting, lowering overall operational expenditure.



12

## AIOps | Drawbacks

- **Implementation Complexity and Cost**
  - AIOps platforms **require significant investment** in **time, resources, and expertise** to **integrate with diverse existing tools** and to customize ML models.
- **Data Quality Dependency**
  - The intelligence of AIOps is **only as good as the data** it consumes.
  - **Poor data quality, missing data, or inconsistent data** can **lead to inaccurate insights or false positives**.
- **Need for Specialized Skills**
  - **Teams require** a blend of **IT Operations knowledge** and **data science/ML expertise** to properly configure, train, and maintain the models.
- **Vendor Lock-in**
  - **Solutions** can be **complex to migrate**, leading to potential dependency on a single vendor's platform.



13

## AIOps | Use Cases

Use Case	How AIOps Helps
<b>Predicting Server/DB Failure</b>	ML analyzes logs & performance metrics to <b>predict failures before they happen</b>
<b>Noise Reduction in Monitoring</b>	Correlates <b>150 alerts</b> → <b>into 1</b> meaningful incident
<b>Automated Incident Ticketing</b>	<b>Creates ServiceNow or Jira ticket</b> with context, RCA, and priority
<b>AI-Suggested Remediation</b>	<b>Recommends probable fix</b> (restart service, clear cache, scale nodes)
<b>Auto-Healing</b>	<b>Executes automated scripts</b> (restart pod, scale cluster, rotate logs)
<b>SLO Protection</b>	<b>Warns when availability or latency trends risk breaching SLOs</b>



14



# Organisational Impact & Continuous Improvement in SRE

15

## Continuous Improvement in SRE

- SRE is not just a set of tools and practices — it's a **cultural and organizational transformation**.
- SRE introduces a **shift from reactive operations to proactive reliability engineering**, helping organizations build **reliable, scalable, and secure systems while continuously improving**.

### 1. Building an SRE Culture in Traditional Ops Teams

- **Traditional Ops teams** are usually reactive.
- **SRE Culture Shift Includes:**
  - Introducing **SLOs, SLIs**, and Error **Budgets** for measuring reliability
  - Encouraging **blameless postmortems** to learn from failures
  - Automating **high-toil tasks** (tickets, deployments, configs, monitoring)
  - **Cross-functional collaboration** between Dev, Ops, QA, Security
  - Shifting from "service support" to "service ownership"

### 2. Scaling Reliability Practices Across Teams

- Once SRE principles take root, the next challenge is **scaling reliability across the organization** — not just in one team or system, but **across all platform**, product, and **teams**.
- **Reliability** becomes **institutionalized**, not dependent on **individuals**.



16

## Scaling SRE in an Enterprise | Use Case

- A large enterprise has multiple application teams with inconsistent reliability practices.
- Some have monitoring, others don't.
- Some use automation, others manually run deployments.
- Production breaches happen frequently.
- **SRE-driven Solution:**
  1. Introduced Service Ownership Model → Teams own their reliability
  2. Created a Standard Reliability Playbook (templates for SLOs, alerts, resilience tests)
  3. Established a Central SRE Enablement Team
  4. Onboarded each app team to:
    - o Define SLOs/SLIs
    - o Implement monitoring and auto-rollback
    - o Use golden CI/CD templates
  5. Introduced Quarterly Reliability Reviews and Error Budget tracking
- **Outcome:**
  - o Incidents reduced
  - o Faster recovery (MTTR improved)
  - o Consistent automation across teams



17

**Use Case**



18

# FinEdge Payments

- **Company:** FinEdge Payments (Online payment processing platform)
- **Domain:** FinTech (High availability, transaction integrity, real-time performance needed)
- **Tech Stack:** AWS, Kubernetes, Java microservices, MySQL, Kafka, Istio, Grafana, Prometheus, BigPanda, PagerDuty
- **Daily Transactions:** ~8 million
- **Reliability Objective:** 99.95% Availability (SLA), Error Budget = 21 minutes/month



19

# Problem Before SRE

Challenge	Impact
Too many alerts (noise & false alarms)	<b>Alert fatigue</b> , missed critical issues
Manual incident resolution	<b>Long MTTR</b> (75 mins avg)
No visibility into service dependencies	<b>Hard to diagnose issues</b>
Reactive troubleshooting	<b>Issue detected after customer impact</b>
Recurring incidents	<b>No automation or RCA intelligence</b>

- **System Availability:** 99.6%
- **Incidents Per Month:** 22
- **MTTR:** 75 mins average



20

## SRE Intervention Strategy

- **Introduced Observability (Not Just Monitoring)**

Tool	Use
Prometheus	Metrics collection
Grafana	Visualization, dashboards
Loki / ELK	Log analytics
Jaeger	Distributed tracing
BigPanda	AI-based event correlation & RCA

- **Enabled Full-stack visibility** — metrics, logs, traces, dependencies.



21

## SRE Intervention Strategy ... continue

- **Implemented SLOs & Error Budgets**

Service	SLO	Error Budget (Monthly)
Transaction API	99.95% availability	21 mins
Payment Gateway	<400 ms latency	5 hrs. slow response budget

- Triggered **SLO-based alerting** — alerts only when customer experience was at risk, reducing noise.



22

## SRE Intervention Strategy ... continue

- Automated Incident Detection & Noise Reduction Using AIOps

Capability	What It Did
Event Correlation	130 alerts grouped into 1 root cause incident
Anomaly Detection	Detected DB latency spike before outage
Predictive Analytics	Forecasted API failure due to memory leak
AI-based RCA	Suggested root cause in 3 min instead of 40

- Integrated BigPanda + PagerDuty for **AI-driven alerting and ticket generation.**



23

## SRE Intervention Strategy ... continue

- Auto-Remediation Enabled for Recurring Incidents

Issue Type	Automated Remediation	Issue Type
Memory spike in API pods	Auto-restarts pods using Kubernetes HPA	Memory spike in API pods
Kafka queue overload	Autoscaling via Lambda trigger	Kafka queue overload
DB connection timeout	Auto-executes script to recycle connections	DB connection timeout

- Reduced MTTR from **75 mins** → **12 mins** (84% faster).



24

## Business Outcome After Applying SRE

Metric	Before SRE	After SRE
Incidents per month	22	6
MTTR (Mean Time to Recovery)	75 mins	12 mins
Alert Noise	320/month	70/month
On-call Fatigue	High	Low
Availability	99.6%	99.94%
Customer Complaints	High	Reduced by 60%



25

## Case Study



26

# Grocery Delivery Startup

- “QuickBasket”, a grocery delivery startup, currently uses **Traditional Ops**.
  - Current Status

Category	Metric / Detail	Value / Description
<b>Availability &amp; Downtime</b>	Average monthly downtime	<b>6 hours</b>
	Current uptime	<b>99.17%</b>
<b>Traffic Pattern</b>	Weekend peak surge	<b>4x normal traffic</b>
	Checkout API P95 latency	<b>650 ms</b>
<b>Performance (Latency)</b>	Expected latency target	<b>&lt; 300 ms</b>
	Ops team toil	<b>65% of total work</b>
<b>Weekly Repetitive Toil Tasks</b>	Manual pod restarts	<b>80 per week</b>
	Manual scale-up actions	<b>40 per week</b>
	Ticket-based approvals	<b>25 per week</b>
	Manual config updates	<b>12 per week</b>
<b>Deployments</b>	Deployment method	<b>100% manual, midnight releases</b>
	Deployment failure rate	<b>18%</b>



27

# Grocery Delivery Startup

- “Q1 — Embrace Risk Using SLOs & Error Budgets
  - Using the data above:
    - Propose **two SLOs** for QuickBasket (Availability + Latency).
    - Calculate the **monthly error budget (in minutes)** for your availability SLO.
  - Explain how the error budget should influence release decisions during weekends (4x traffic surge).
  - Explain how this replaces the “zero downtime” mindset of Traditional Ops.
- Q2 — Reduce Toil (Automation Required)
  - Using the toil numbers given:
  - Identify **three major toil items** from the list.
  - Calculate **total weekly toil actions** (sum them).
  - Recommend which tasks should be automated first and **why**, using numbers.
  - Explain how this reduction in toil will improve:
    - Reliability
    - MTTR
    - Team productivity
- Q3 — Engineering-Focused Operations (Blameless, Monitoring, Coding Ops)
  - Based on the scenario:
  - Describe **how blameless postmortems** would improve reliability compared to current blame-driven ops.
  - Propose at least **3 SLIs (latency, availability, freshness, etc.)** to rebuild monitoring.
  - Explain how adopting **Infrastructure as Code / Monitoring as Code** will fix the current repetitive manual tasks.
  - Explain how Dev + Ops collaboration will reduce the **18% deployment failure rate**.



28

# Business Case for Adopting Site Reliability Engineering (SRE)



29



## Current Environment Overview



30

## Current Environment

- The organization operates a **large on-premises infrastructure estate** with multiple technology platforms, databases, private cloud systems, middleware, and storage environments.
- Each domain uses **multiple monitoring tools**, leading to tool fragmentation and operational complexity.

Infra Component	Tech Stack	Monitoring Tools
Unix / Windows	Platforms	SCOM, Datadog, Compass, Nagios, Truesight, ServerSiter, iMon, EMAT
Oracle, MySQL, Hadoop, MongoDB	Database & Big Data	OEM, DataDog, Truesight, Cloudera, Splunk
NAS, SAN, Rubrik	Storage & Backup	Hitachi OpsCenter, Rubrik CDM/Netbackup, Pure1, RSE
VMware Hosting	Private Cloud	vROPS (vRealize Aria Operations)
IBM MQ, SWIFT	Middleware	TruSight, Datadog

- This **fragmented tooling environment** increases operational cost, complexity, and hampers efficient decision-making, particularly during critical incidents.



31

## Business Challenges

- The current operational model faces the following challenges:
  1. Fragmented Monitoring & Lack of Observability
  2. Difficult to Consolidate Monitoring Tools
  3. Traditional Ops Mindset
  4. Lack of Standardized Toil Identification
  5. Limited Automation & IaC Adoption
  6. No Business-Aligned SLOs Defined
  7. Capacity Management Requires Improvement



32

# Transitioning from Traditional ITOps to SRE Recommendations

AHMAD

33

## Landscape

- Your infrastructure spans multiple layers (*Platform, DB, Storage, Cloud, Middleware*) and uses **multiple isolated monitoring tools**, with focus mainly on **reactive monitoring** and **availability management**.

Current State	Target SRE State
Tool-centric monitoring	<b>Service-centric observability</b>
Reactive incident handling	<b>Proactive + predictive reliability management</b>
Ticket-driven Ops	<b>Automation</b> , IaC, AI-driven insights
High manual toil	<b>Reduced Toil</b> , Continuous improvement
Infra availability focus	<b>Measure business reliability</b> via SLOs, SLIs

AHMAD

34

17



35

## Shift Mindset

- The **shift** from a traditional operational model (*where the primary goal is often avoiding failure and handling manual tasks*) to an **SRE mindset** (*focused on engineering reliability and automation*) requires a structured change management approach.
- **Mindset Shift Activities**
  - Conduct **SRE workshops**: SLO, Error Budget, Blameless Postmortems, Automation principles.
  - Create **Reliability Champions** within each platform team.
    - Reliability Champions are **nominated engineers within each infrastructure/platform team who act as the evangelists of SRE practices**, guiding their team in adopting reliability, automation, observability, and error budgeting principles.
  - Adopt **Production Readiness Reviews (PRR)** before changes.
    - A **Production Readiness Review (PRR)** is a **structured assessment** conducted **before deploying any major change**, application, system, or infrastructure component to production, to ensure security, reliability, observability, scalability, and automation readiness.
    - It acts as an **SRE gate** to verify that the system meets reliability standards before release.



36

## Shift Mindset ... continue

Traditional Mindset	SRE Mindset	Implementation Strategy
"Keep the Lights On"	"Automate the Operations"	<ul style="list-style-type: none"> <li>Introduce the <b>Error Budget</b> concept.</li> <li>When the service is reliable enough (<i>error budget is high</i>), the team can ship features; when the error budget is depleted, they stop and focus on reliability work (<b>TOIL reduction, bug fixes</b>).</li> </ul>
Manual Escalation/Fixes	Eliminate Toil	<ul style="list-style-type: none"> <li>Mandate that SREs spend <b>50% of their time</b> on engineering work (automation, new features, design).</li> <li>Any operational work exceeding 50% must be automated or eliminated.</li> </ul>
Blame-Oriented Reviews	Blameless Postmortems	<ul style="list-style-type: none"> <li>Establish a culture where incidents are viewed as <b>systemic failures</b>, not personal ones.</li> <li>Focus on <b>what</b> happened, <b>why</b> it happened, and <b>how</b> to prevent recurrence, documenting actions in a <b>postmortem document</b>.</li> </ul>



37

Tool Consolidation

38

19

## Tool Consolidation

- Your current tool landscape is highly fragmented (*SCOM, Nagios, TruSight, Datadog, etc.*), which creates complexity and slows down incident response.
- **Vision Tool**
  - **Datadog** is already present and offers strong capabilities across **Metrics, Logs, and Traces (the three pillars of Observability)**, making it a viable candidate for consolidation.
- **Strategy**
  - Implement a phased approach using **OpenTelemetry (OTEL)**.
    - **OpenTelemetry**
      - Use **OTEL agents and collectors** to standardize the data format *before* it gets sent to a monitoring backend.
      - This allows you to **collect data** once and send it to your **existing tools and the new consolidated tool** (**Datadog**) simultaneously.
      - This **de-risks the migration**.



39

## Tool Consolidation ... continue

- **Phased Migration**
  - **Migrate one component** (e.g., *your VMware Private Cloud*) completely to **Datadog/OTEL**, prove the new system's value, and then **decommission** the **old tools** (*vROPS, TrueSight*) for that component.
- **Data Correlation**
  - A unified tool like **Datadog** will allow for automatic **data correlation across platforms**, databases, and middleware, which is critical for incident root cause analysis.



40




41

## Eliminating TOIL

- **TOIL** (work that is manual, repetitive, automatable, tactical, and lacks enduring value) is a key SRE focus area.
- **Better Ways to Identify TOIL:**
  - **Time Tracking**
    - Mandate SREs and **Ops engineers** to **log their time** for **every operational task** (e.g., "manual server patching," "routine database backup verification").
    - **Any task** that takes **>5 minutes** and is **repeated >10 times a week** is a prime **TOIL candidate**.
  - **Alert Review**
    - **Identify the top 10** most **frequent alerts**.
    - If an **alert** is acknowledged or **resolved without** a **code change** or system configuration change, the resolution process is likely TOIL and **should be automated**.
  - **Postmortem Analysis**
    - Every **incident postmortem** should **explicitly identify** if any part of the resolution was **TOIL** and propose an **automation** task to eliminate it.
  - **Engineer Survey**
    - **Conduct** a simple, regular (e.g., monthly) **survey** asking: "What is the most annoying, repetitive task you did this month?"
- **Automation Tools**
  - **Ansible** or **Puppet** are excellent for automating OS/Platform TOIL (e.g., patching, configuration drift).
  - **Python scripting** is key for automating database/middleware tasks.



42




43

## Observability with Tracing

- You are currently monitoring Metrics (e.g., CPU, Memory) and Logs (e.g., application error messages).
- To **achieve end-to-end observability**, you **need Traces**.
- Implementation Example: **Tracing Middleware/Database Calls**
  - **Goal**
    - **Track a transaction's journey** from the platform (*Unix/Windows*) through the middleware (*IBM MQ*) to the database (*Oracle/MySQL*).
  - **Tool**
    - **Implement OpenTelemetry (OTEL)** or **use Datadog APM (Application Performance Monitoring)**.
- **Benefit**
  - If a **transaction takes 10 seconds**, a trace will **show exactly how much time was spent** in the **application**, how long the **message** sat in the **MQ queue**, and **how long** the **database query took** (e.g., \$9\$ seconds spent in the database).
  - This pinpoints the **performance bottleneck immediately**.



44



45

## Automation

- IaC is fundamental to SRE, reducing TOIL and ensuring reproducibility.
- Recommended **IaC Tools**:
  - **Platform (OS/VMware)**
    - **Terraform** for **provisioning** and **managing** your VMware Cloud infrastructure (VMs, networks, storage).
    - This is your primary IaC tool.
  - **Configuration Management**
    - **Ansible** or **Puppet** for **configuring** the OS (*Unix/Windows*), **installing** necessary packages, and hardening the platforms.



46

## Automation ... continue

- Automation Opportunities
  - Self-Healing/Auto-Remediation
    - Automate responses to common alerts (e.g., if a server's disk utilization exceeds 90%, automatically run a script to clear temporary files and alert the team).
  - Environment Provisioning
    - Automate the end-to-end build of a new environment (Platform, Middleware, Database) using Terraform and Ansible/Puppet.
  - AIOps Integration
    - Use your new AIOps tool (BigPanda) to ingest alerts from all your remaining monitoring tools and automatically trigger these Ansible/Puppet automation playbooks for remediation.



47



48

# BigPanda

- **BigPanda** acts as an **Event Correlation Engine** and **Noise Reduction tool**.
- **Role of BigPanda:**
  - Ingestion**
    - **Ingest alerts** and **data** from all your **disparate tools** (*SCOM, Nagios, vROPS, Datadog*).
  - Correlation**
    - **Use Machine Learning** to **group related alerts** into a single **Incident**.
    - For example, a single network switch failure might trigger 50 alerts in Nagios, vROPS, and SCOM; BigPanda should consolidate this into **one incident**.
  - Automation Trigger**
    - The consolidated incident in BigPanda should be used as the **single source of truth** to:
      - **Automatically open a ticket** in your **ITSM** system (e.g., ServiceNow).
      - Automatically trigger the necessary **Ansible/Puppet automation playbook** for **auto-remediation**.



49

Service Levels

50

## SLOs

- SRE shifts focus from **component health** to **service health**.
- You need to translate your infrastructure's health into **metrics** that **matter** to the **business**.
- **Infrastructure SLOs based on Business Impact:**

Component	SLO Metric (SLI)	SLO Target	Business Implication
Private Cloud Compute	VM Provisioning Latency	95% of VMs provisioned in <b>&lt;5 minutes</b>	Time to Market for new services/fixes
Storage & Backup	Backup Success Rate	<b>&gt;=99%</b> success rate on nightly backups	Data Loss Risk (RPO)
Database	Query Latency (P99)	<b>99%</b> of critical read queries complete in <b>&lt;100 ms</b>	Application Performance and User Experience
Middleware (IBM MQ)	Message Queue Depth/Latency	Message queue latency is <b>&lt;2</b> seconds for <b>99.9%</b> of transactions	Transaction Throughput and Reliability

- **Error Budget**
  - Once SLOs are set, the Error Budget is **1 - SLO** (e.g.,  $1 - 99.9\% = 0.1\%$  acceptable downtime/error).



51

## Capacity Management



52

# Capacity Management

- In SRE, capacity management moves from a reactive, annual budget process to a **proactive, data-driven forecasting** model tied to utilization and performance.
  - **Leverage Observability Data**
    - Use your consolidated tool (*Datadog*) to establish **baseline utilization** (*CPU, Memory, Disk, Network I/O*) for **all critical infrastructure components**.
  - **Saturation SLIs**
    - **Define Saturation SLIs** (e.g., *CPU Utilization P99, Queue Length*) and set warning **alerts** when they consistently **exceed a threshold** (e.g., 70%).
  - **Growth Forecasting**
    - **Correlate business volume metrics** (e.g., *number of user transactions, number of active database connections*) with infrastructure resource consumption.
    - **Use this correlation to forecast** when you will hit the **70% saturation threshold** and initiate the **capacity procurement** process before a capacity-related incident occurs.
  - **Virtualization Optimization**
    - Use **vROPS** data to identify and **reclaim zombie VMs** (*powered-off or underutilized VMs*) and right-size over-provisioned VMs to maximize resource efficiency.



53

## Q & A

Any concepts still unclear?

---

**Thank you for attending**



54