

## SRE Automation Case Study

### Stock Exchange Trading Platform

#### Industry Context

A national stock exchange runs a **high-frequency trading (HFT) platform**. Peak traffic occurs during:

- Market open (9:00–9:30 AM)
- Major announcements
- Sudden buy/sell surges

They must ensure:

- Millisecond latency
- Zero downtime during trading hours
- Fast recovery
- Automated reliability operations

#### Case Study Scenario

**ExchangeX**, a major stock exchange, operates a trading engine that processes **50 million transactions/day**.

#### Current Challenges

1. **Manual deployment approval**, leading to delays.
2. **Frequent CPU spikes** on matching engine (buy/ sell) nodes.
3. **Order processing queue delays**, not detected until traders complain.
4. **Manual rollbacks** that take 20–30 minutes.
5. **Incidents are manually triaged**, slowing MTTR.
6. **Inconsistent logs**, making root-cause analysis difficult.

#### Business Goals

- Achieve **99.99% availability** during trading hours.
- Reduce **MTTR from 25 minutes to <5 minutes**.
- Automate **scaling, deployment, rollbacks**, and **incident detection**.
- Reduce **manual operations by 70%** within 3 months.

## **Questions**

- Q1.** What SRE Automation should be introduced to reduce manual deployment delays?
- Q2.** How can auto-scaling be implemented for a trading engine that cannot go down?
- Q3.** What automation can detect order queue delays before traders report issues?
- Q4.** How can automated rollbacks be implemented for critical trading services?
- Q5.** What SRE automation can reduce MTTR during high-severity incidents?
- Q6.** What logging and monitoring automation strategy would best fit a stock exchange platform?
- Q7.** How can SREs use resiliency testing safely in a financial trading system?
- Q8.** What error budget policies should be set for a 99.99% SLA trading system?

## Model Answer

<b>Q No.</b>	<b>Question</b>	<b>Model Answer</b>
A1	What SRE Automation should be introduced to reduce manual deployment delays?	<b>Automated CI/CD + GitOps + Canary Deployment:</b> Trigger builds automatically from Git, auto-deploy to staging, then canary deploy to 1% of nodes; auto-rollback if latency/error rate spikes. <b>Removes manual approvals and accelerates safe releases.</b>
A2	How to implement auto-scaling for a trading engine that cannot go down?	<b>Predictive + Layered Scaling:</b> Pre-scale trading systems before peak hours using historical traffic; gateway layers use horizontal auto-scaling; core matching engines stay static but redundant; use blue-green for safe replacement. <b>Ensures zero-downtime scale adjustments.</b>
A3	What automation can detect order queue delays early?	<b>SLO-based Monitoring Automation:</b> Track queue depth, order latency, and processing rate; set automated alerts; auto-run actions such as scaling gateways, restarting slow consumers, or notifying teams via Slack. <b>Detects issues before traders complain.</b>
A4	How to automate rollbacks for critical trading services?	<b>Automated Health Check Rollbacks + Feature Flags:</b> CD pipeline checks latency/error metrics; if 3 checks fail, auto-rollback triggers; maintain versioned images; feature flags instantly disable problematic logic. <b>Rollback completes within &lt;60 seconds.</b>
A5	What SRE automation can reduce MTTR in high-severity incidents?	<b>Runbook Automation + AIOps + ChatOps:</b> Automatically diagnose issues (CPU spikes, latency, queue backlog); auto-run common recovery actions; incident bot posts logs, root cause hints, and suggests runbooks. <b>MTTR drops from 25 minutes → &lt;5 minutes.</b>
A6	What logging/monitoring automation suits a stock exchange?	<b>Centralized Structured Logging + Auto-Anomaly Detection:</b> Use JSON logs, OpenTelemetry, and ELK/Loki pipeline; automate correlations (order ID ↔ error ↔ latency); auto-trigger alerts on anomalies; enforce automated log retention. <b>Speeds RCA and ensures compliance.</b>
A7	How can SREs use resiliency testing safely in a financial trading system?	<b>Run tests only during non-trading hours</b> Perform tests in staging or pre-production, not live markets; test small failures like: one server slowing down, network delay, or a service restart Ensure backup systems take over smoothly Always limit the blast radius (only test one component at a time). This helps the <b>exchange verify that trading continues even if a part of the system fails</b> —without risking real-money trades.
A8	What error budget policy fits 99.99% SLA?	<b>52 min/year downtime → Strict Budget Policy:</b> Only 40 minutes considered usable; freeze deployments after 70% consumption; allow only canary/feature-flag releases when 80–100% budget is used. <b>Protects stability during trading hours.</b>