

## Case Study

"**QuickBasket**", a grocery delivery startup, currently uses **Traditional Ops**.

### **Current Metrics**

- Average monthly downtime: **6 hours**
- Current uptime: **99.17%**
- Peak order surge: **4x traffic every weekend**
- Checkout API latency:
  - P95 latency: **650 ms**
  - Expected: **< 300 ms**
- Ops team toil: **65%** of total work
- Weekly repetitive tasks:
  - 80 manual pod restarts
  - 40 manual scale-up actions
  - 25 ticket-based approvals
  - 12 manual config updates
- Deployment:
  - 100% manual, done at midnight
  - Failure rate: **18%**

The CTO wants to shift to **SRE model** using **three principles**.

### **Task**

#### **Q1 — Embrace Risk Using SLOs & Error Budgets**

Using the data above:

1. Propose **two SLOs** for QuickBasket (Availability + Latency).
2. Calculate the **monthly error budget (in minutes)** for your availability SLO.
3. Explain how the error budget should influence release decisions during weekends (4x traffic surge).
4. Explain how this replaces the "zero downtime" mindset of Traditional Ops.

## **Q2 – Reduce Toil (Automation Required)**

Using the toil numbers given:

1. Identify **three major toil items** from the list.
2. Calculate **total weekly toil actions** (sum them).
3. Recommend which tasks should be automated first and **why**, using numbers.
4. Explain how this reduction in toil will improve:
  - o Reliability
  - o MTTR
  - o Team productivity

## **Q3 – Engineering-Focused Operations (Blameless, Monitoring, Coding Ops)**

Based on the scenario:

1. Describe **how blameless postmortems** would improve reliability compared to current blame-driven ops.
2. Propose at least **3 SLIs (latency, availability, freshness, etc.)** to rebuild monitoring.
3. Explain how adopting **Infrastructure as Code / Monitoring as Code** will fix the current repetitive manual tasks.
4. Explain how Dev + Ops collaboration will reduce the **18% deployment failure rate**.

## Model Answer

### Q1: Embrace Risk Using SLOs & Error Budgets

#### 1. Proposed SLOs

- **Availability SLO:** 99.5% monthly
- **Checkout API Latency SLO:** P95 < 300 ms

#### 2. Monthly Error Budget

For **99.5% SLO:**

- Month = 30 days = 43,200 mins
- Allowed downtime = 0.5%
- **Error Budget = 216 minutes per month**

#### 3. How Error Budget Influences Weekend Releases

Since traffic spikes **4x every weekend**, releases should follow rules:

- If more than **150 mins** of error budget already consumed → **freeze releases**
- If less than **50 mins** consumed → allow releases with rollback alarms
- If latency SLO is violated ( $P95 > 300 \text{ ms}$ ) → reduce release speed and fix hot paths

This transforms decisions into **data-driven governance** rather than gut feeling.

#### 4. Replacing Zero Downtime Mindset

Traditional Ops: "No downtime allowed → delay releases → slow innovation."

SRE: "Downtime is allowed within budget → controlled risk → faster delivery + predictable reliability."

### Q2: Reduce Toil

#### 1. Three Major Toil Items

- 80 manual pod restarts
- 40 manual scale-ups
- 25 ticket-based approvals

#### 2. Total Weekly Toil Actions

$$80 + 40 + 25 + 12 = \mathbf{157 \text{ repetitive actions per week}}$$

### 3. What to Automate First & Why

- **Pod restarts (80/week)** → Highest volume & easily scriptable
- **Scale-up actions (40/week)** → Causes latency drops during spikes
- **Ticket approvals (25/week)** → Causes deployment delays and long MTTR

Automating these removes **145/157 = 92%** of repetitive actions.

### 4. Reliability Benefits

- **Lower MTTR:** Auto-healing replaces manual restarts
- **Better Availability:** Auto-scaling handles spikes instantly
- **Improved Productivity:** 65% toil reduced → engineers focus on redesigning slow services
- **Fewer Errors:** No human mis-clicks / wrong config edits

## Q3: Engineering-Focused Operations

### 1. Blameless Postmortems

Current method:

- “Who caused it?” → engineers hide mistakes → repeated issues

SRE method:

- Ask “What failed?” → focus on systemic causes
- Creates shared learning → eliminates recurring incidents
- Allows bold changes → improves culture

### 2. Proposed SLIs

- **API latency (P95)**
- **Availability (successful requests / total)**
- **Order freshness time** (time from cart → backend processing)
- **Queue lag** for inventory updates

### 3. Monitoring & Infrastructure as Code

- “Manual config updates (12 per week)” replaced with Git-based IaC
- Dashboard + alert rules become versioned code
- Alerts tied to SLOs → no more noisy thresholds
- Faster rollbacks using Git history

#### **4. Reduce 18% Deployment Failure Rate**

- Use CI/CD pipelines with automated tests
- Shift-left reliability checks
- SRE + Dev collaboration to define SLO-based gating
- Gradual rollouts (blue/green, canary)

Outcome: Deployment failure rate drops from **18%** → < **5%** over time.